

CARIAD

Master Thesis: AI Usage in CI/CD/CT Pipelines for Compute Platforms in Automotives

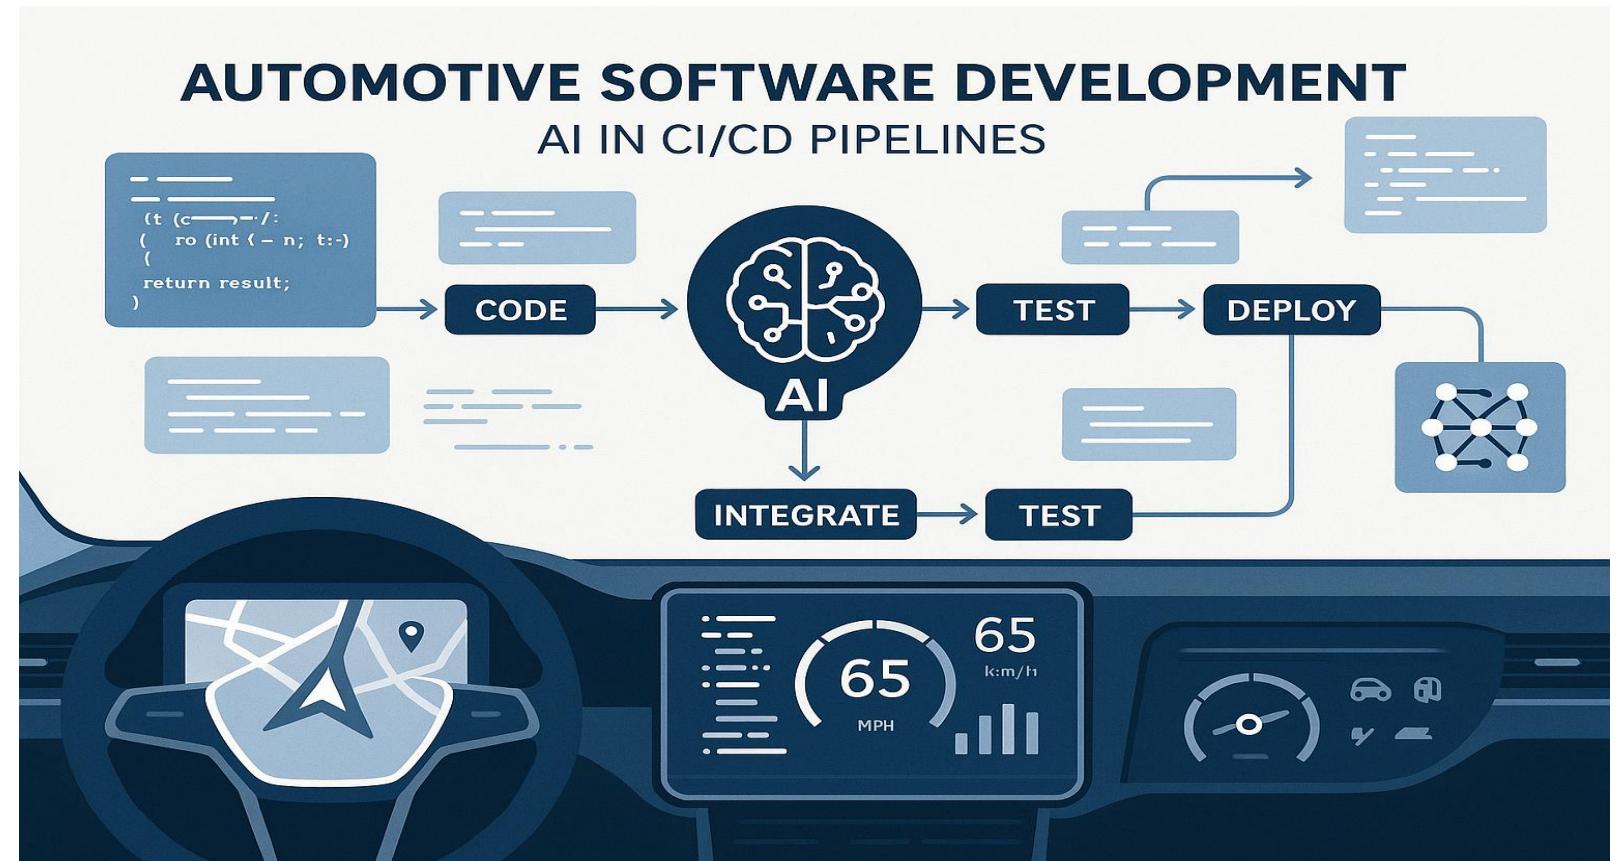


We transform automotive mobility

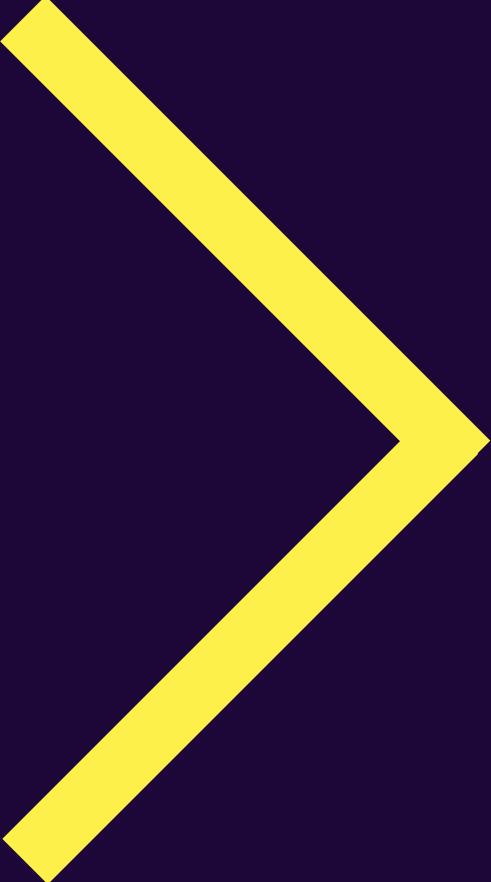
CARIAD
A VOLKSWAGEN GROUP COMPANY

Agenda

- // Introduction
- // Method
- // Result
- /7 Status
- // Next Steps

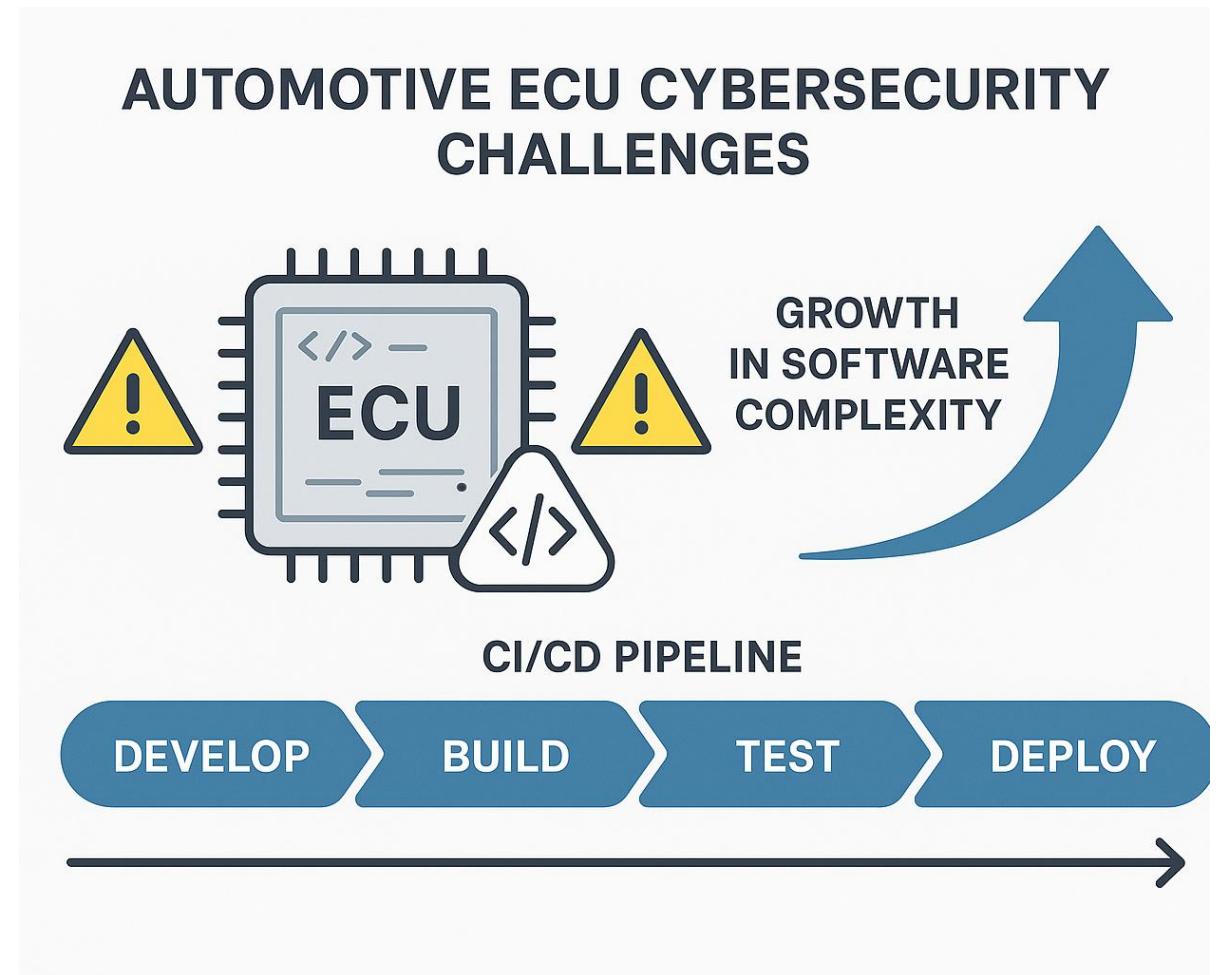


Introduction



Introduction

- **Rapidly growing software complexity & shorter release cycles**
- **Automotive ECUs are safety-critical, hence zero tolerance**
- **Traditional security tests cannot keep pace with CI/CD demand**



Problem Statement



- Current white-box fuzzing & testing are *manual or slow to scale*



- Vulnerabilities may slip through nightly CI due to time limits



- Need an AI-guided approach integrated into CI/CD/CT to



- boost path coverage



- reduce manual fuzz test case creation



- auto-generate actionable test artifacts

Research Objectives

Technique Design

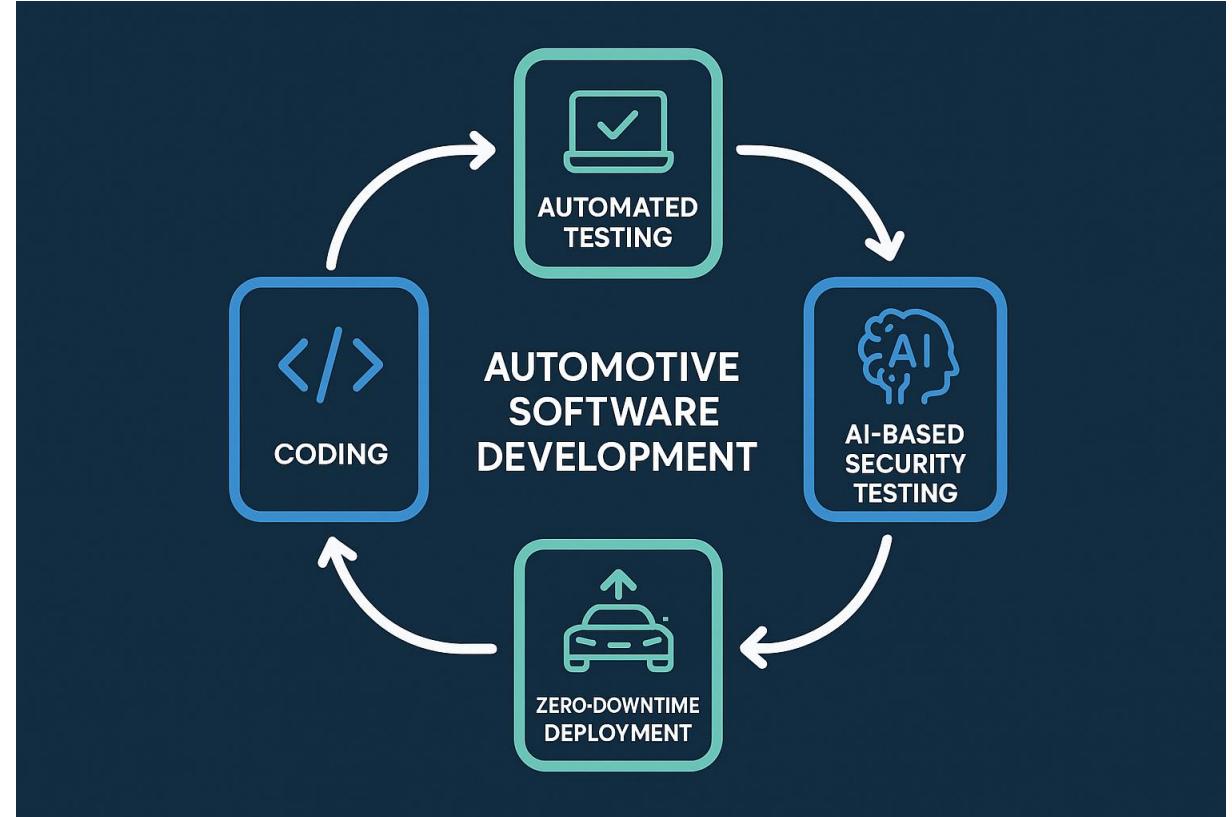
- **AI-assisted white-box fuzzing for automotive targets**

Pipeline Integration

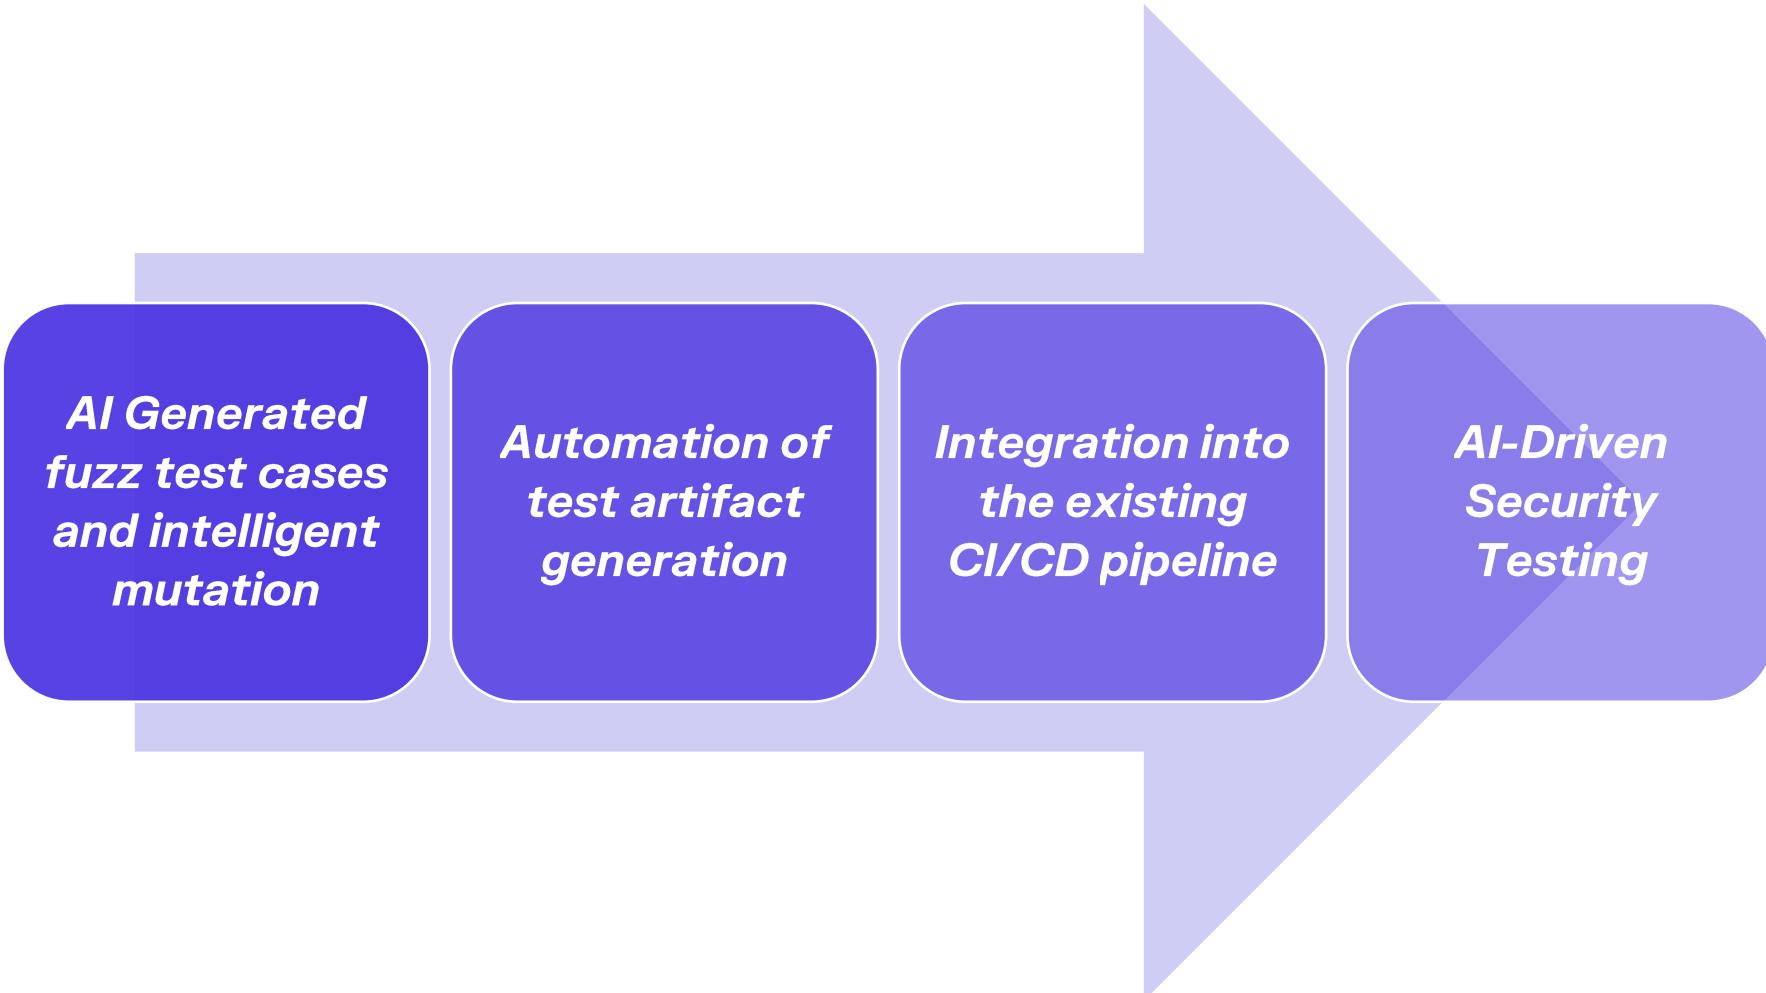
- **Embed continuous fuzzing into existing CI/CD/CT**

Artifact & Impact Automation

- **Auto-generate test cases, reports, quality matrix**
- **Measure coverage, MTTV, and CI latency vs. baseline**



Expected Outcomes



*AI Generated
fuzz test cases
and intelligent
mutation*

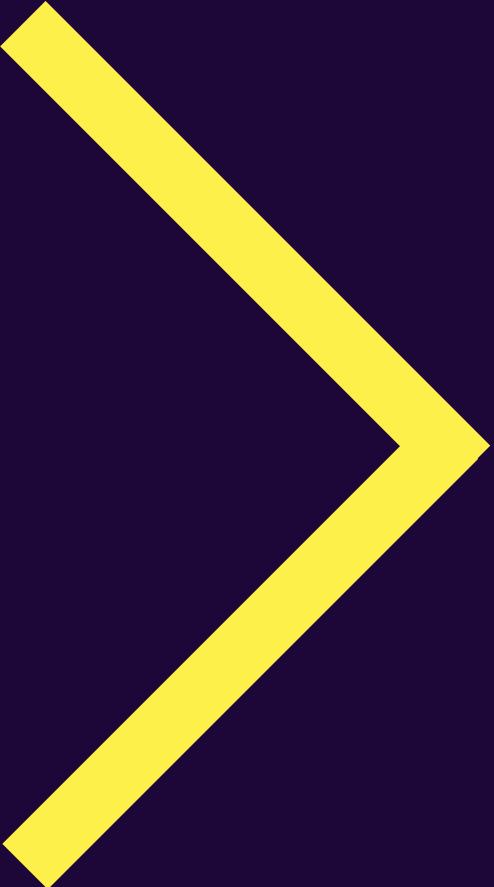
*Automation of
test artifact
generation*

*Integration into
the existing
CI/CD pipeline*

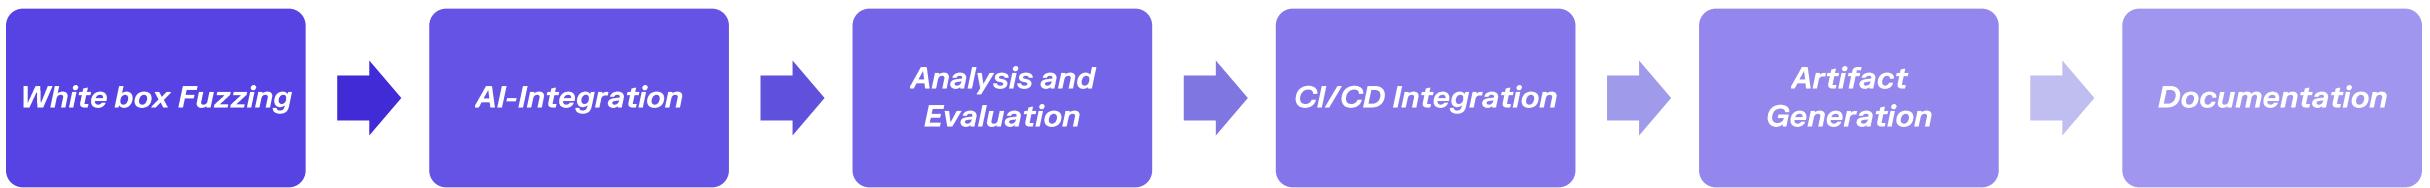
*AI-Driven
Security
Testing*

May

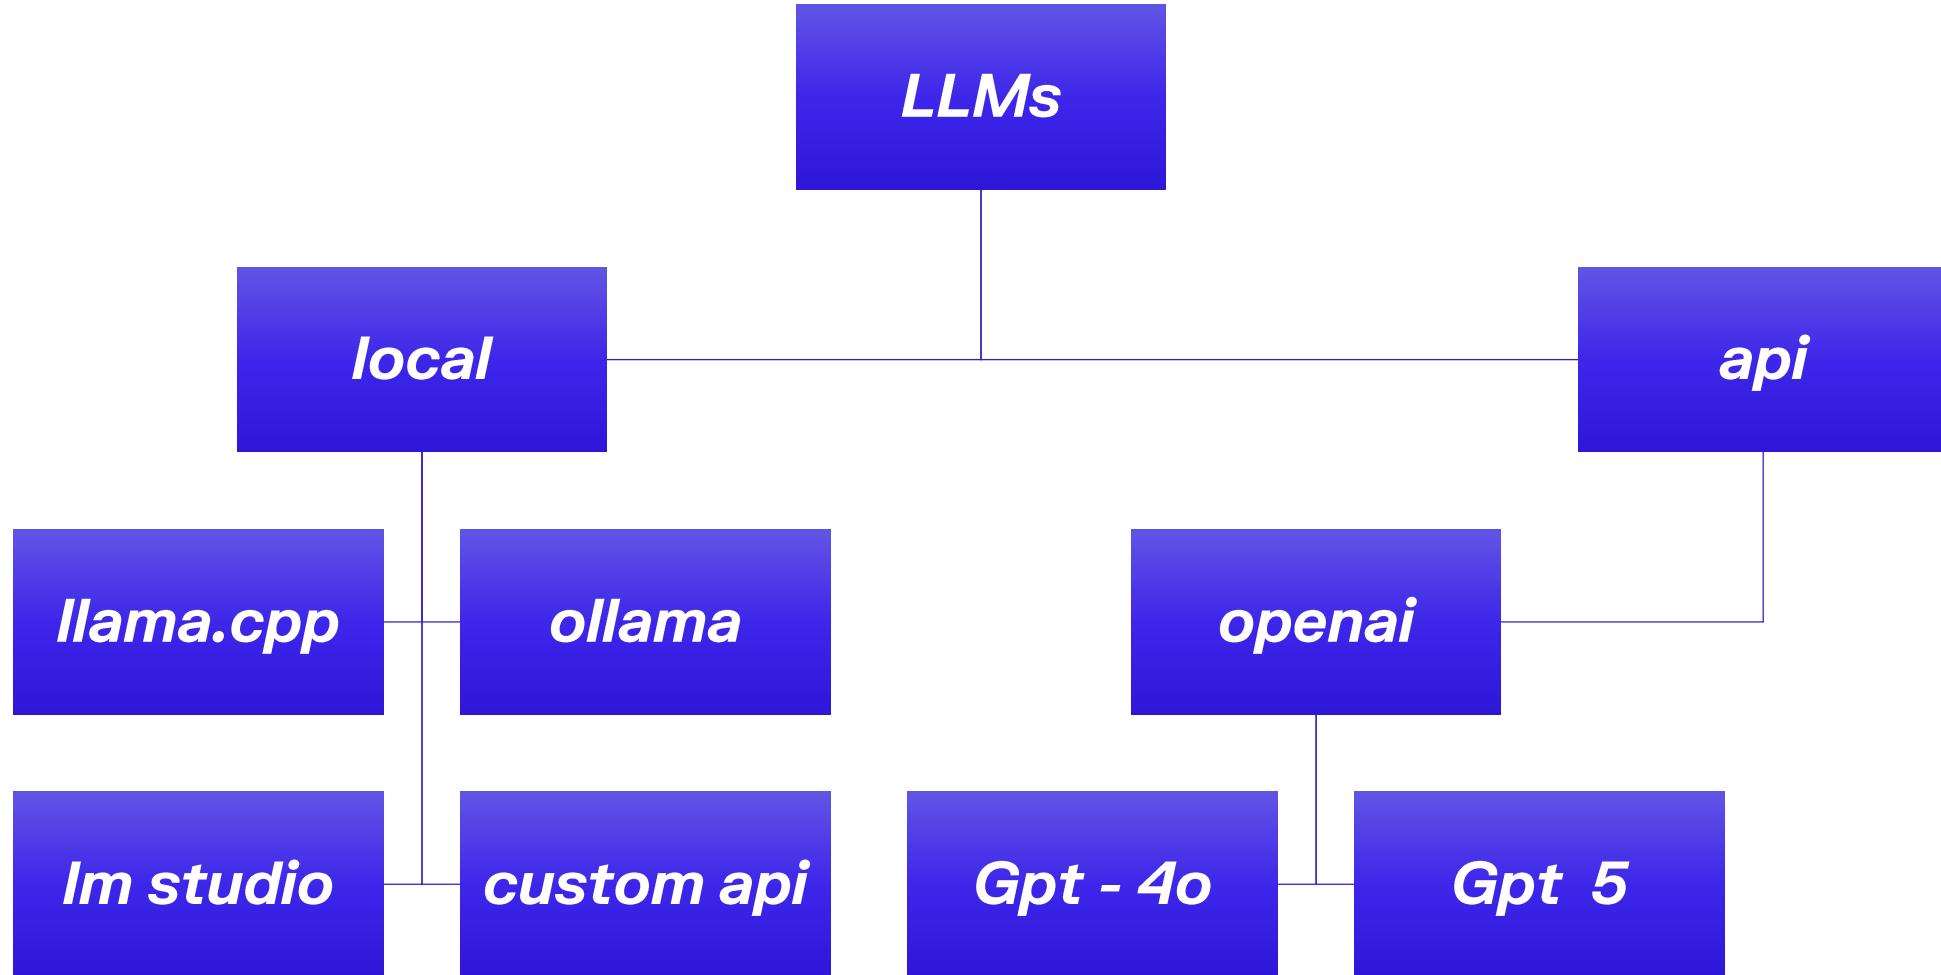
Method



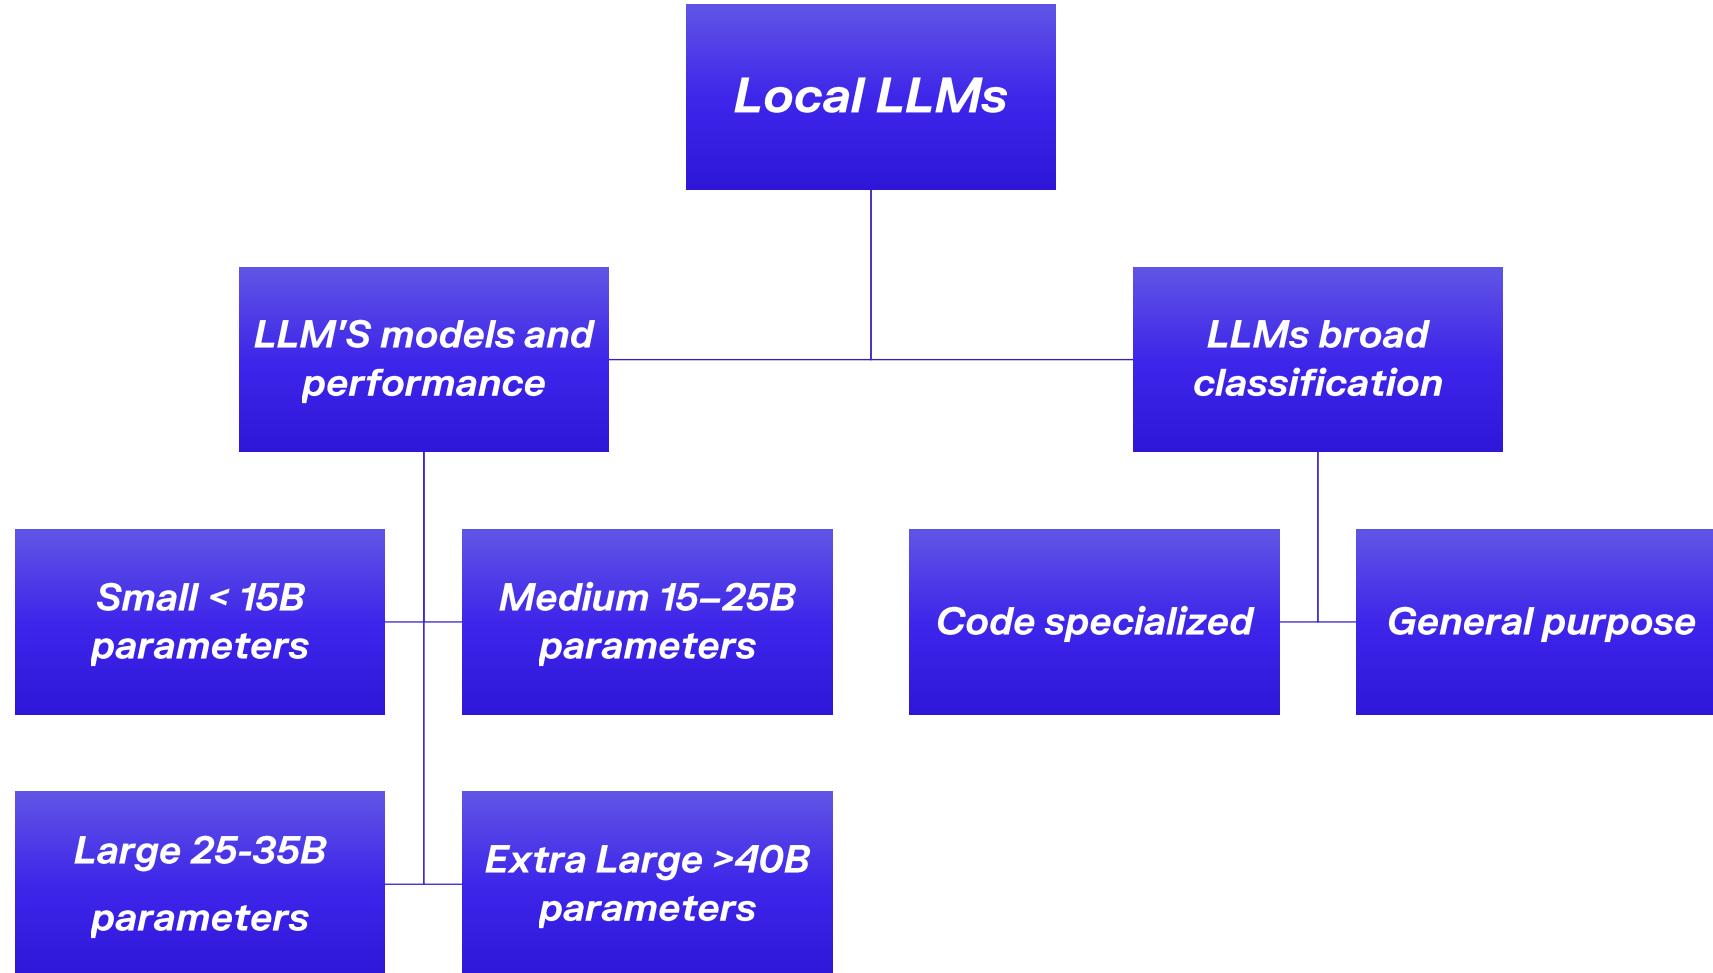
Overview



Approach



Local LLMs



Models

Small (<15B parameters)

- ***phi4:14b (14B)***
- ***llama3:latest (8B)***
- ***qwen2.5-coder:1.5b (1.54B)***
- ***qwen2.5-coder:7b (7B)***
- ***qwen2.5-coder:14b (14B)***

Medium (15–25B parameters)

- ***magistral:24b (24B)***
- ***devstral:latest (24B)***
- ***starcoder2:15b (15B)***

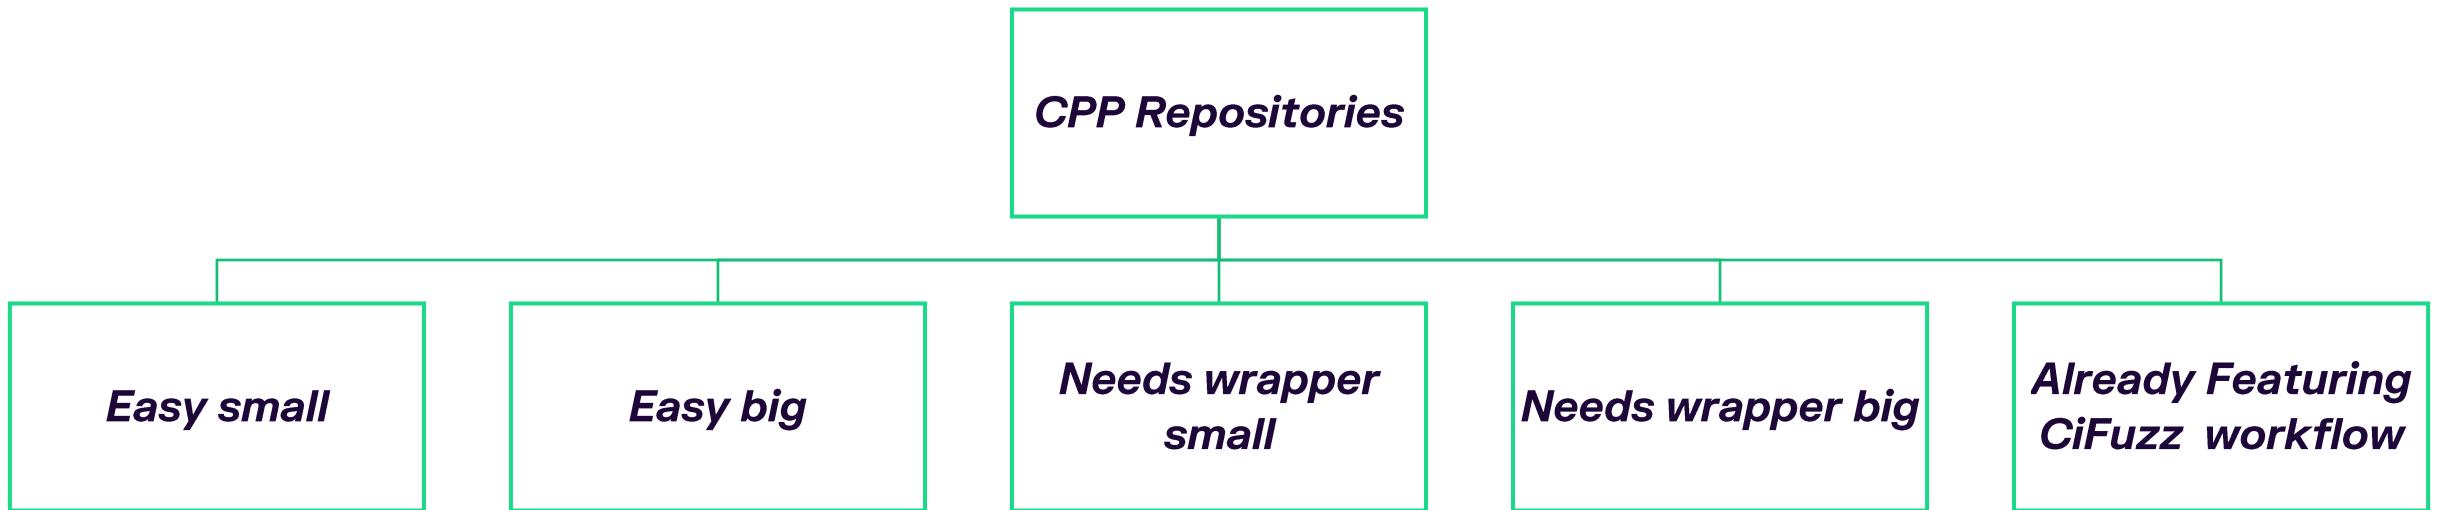
Large (25-35B parameters)

- ***gemma3:27b (27.4B)***
- ***qwen3:32b (32.8B)***
- ***deepseek-r1:32b (32.8B)***
- ***qwen2.5-coder:32b (32.8B)***
- ***deepseek-coder:33b (33B)***
- ***wizardcoder:33b (33B)***
- ***codellama:34b-instruct (34B)***
- ***yi:34b (34B)***

Extra Large (> 40B parameters)

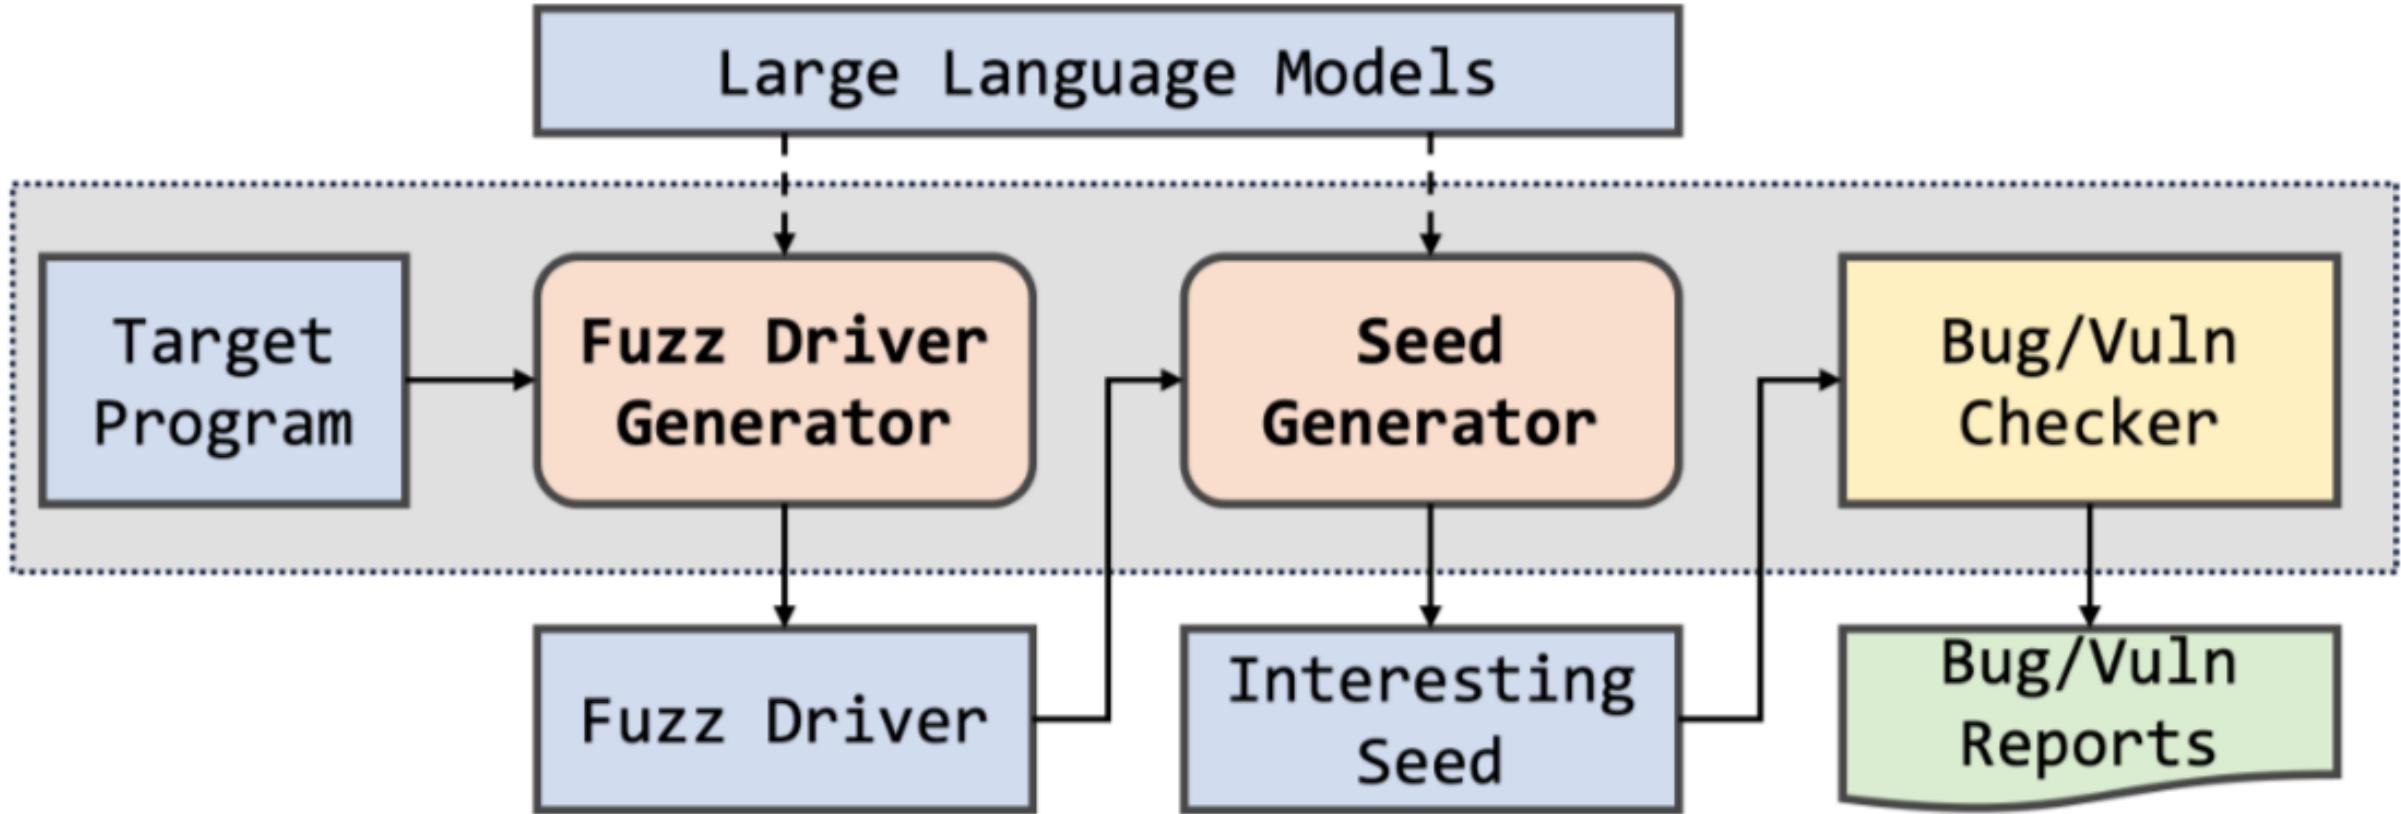
- ***mixtral:latest (46.7B)***

Repositories



repositories

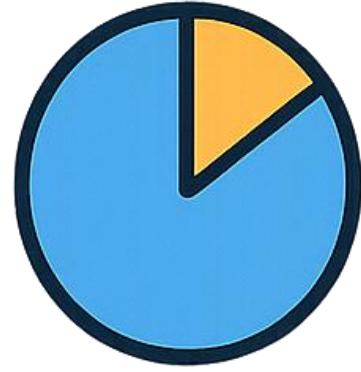
Working



Evaluation Framework



Bugs found



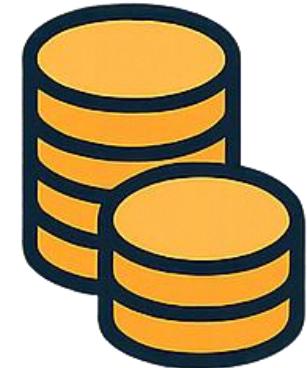
**Code
coverage %**



**Fuzz test
generation**



**Generated
code quality**



**Tokens
used**

Finetuning local llms

Qwen 2.5 coder 32b instruct full model 60gb

LoRA (Low-Rank Adaptation)

- **LoRA rank (r): 16 – controls adaptation size**
- **LoRA alpha: 32 – scaling for adaptation**
- **Dropout: 0.1 – prevents overfitting**
- **Target modules: q_proj , v_proj , etc. – efficient fine-tuning**
- **Device: auto – runs on best available hardware**
- **Dtype: float16 – faster, less memory**
- **Efficient model loading and saving (safetensors)**

Code examples for training

– *Jsonl format*

{"prompt": "Generate a comprehensive parser fuzzer for structured data\\n\\nRequirements:\\n- Maximize code coverage and edge case detection\\n- Generate intelligent, targeted inputs\\n- Optimize for vulnerability discovery speed\\n- Include comprehensive error handling,"}

Code examples for training

Qwen 2.5 coder 32b

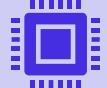
"completion":

```
"void train_timing_attack(const uint8_t* data, size_t size) {\n    if (size < 20) return;\n\n    const char* secret = \"supersecretpassword123\";\n    std::string user_input(reinterpret_cast<const char*>(data),\n                           std::min(size, strlen(secret)));\n\n    auto start =\n        std::chrono::high_resolution_clock::now();\n\n    // VULNERABLE: Early return creates timing difference\n    bool matches = true;\n    for (size_t i = 0; i < strlen(secret) && i < user_input.length(); i++) {\n        if (secret[i] != user_input[i]) {\n            matches = false;\n            break; // Early return reveals timing information\n        }\n        // Simulate processing time\n    }\n\n    auto end =\n        std::chrono::high_resolution_clock::now();\n    auto duration =\n        std::chrono::duration_cast<std::chrono::microseconds>(end - start);\n\n    printf(\"Password check took %ld microseconds, match: %s\\n\", duration.count(), matches ? \"true\" : \"false\");\n}
```

Finetuning



32b – needs more than 60gb ram



14b – (32 – 37gb) ram



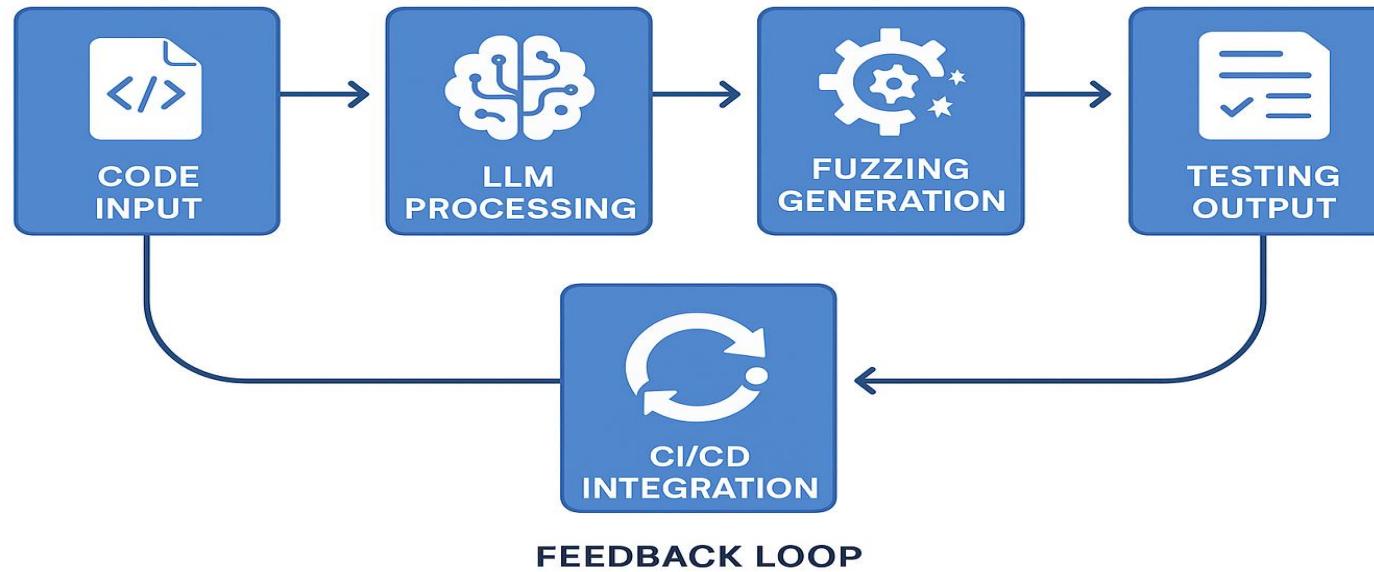
7b – 24gb ram without any background process



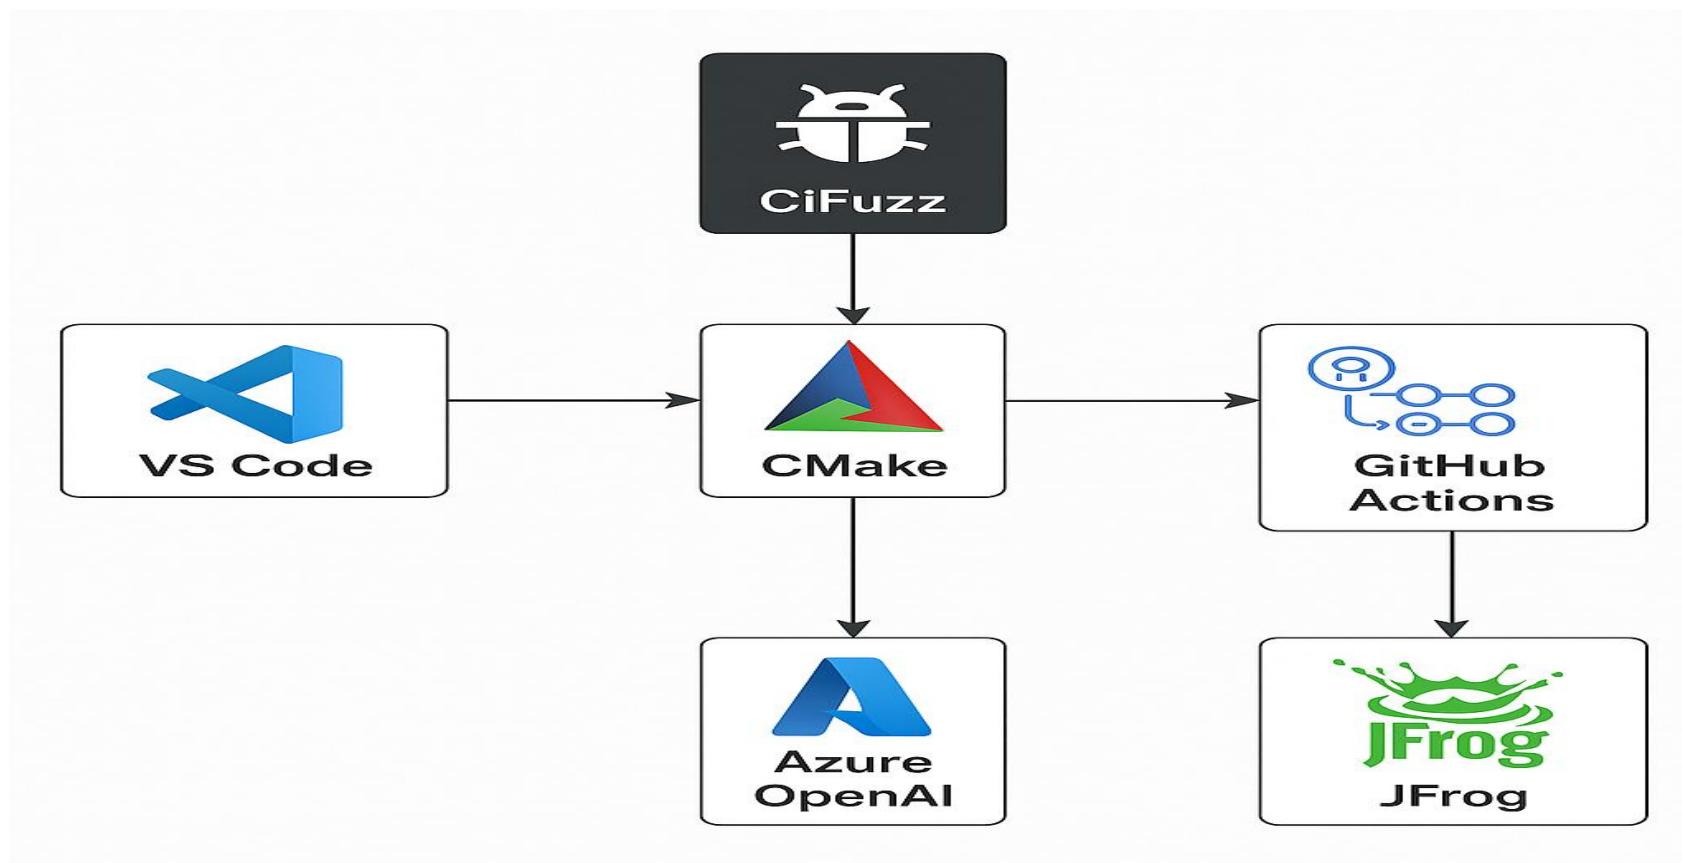
1.5b – 14gb ram

CI/CD Process flow

Process Flow Diagram for Automotive AI Fuzzing

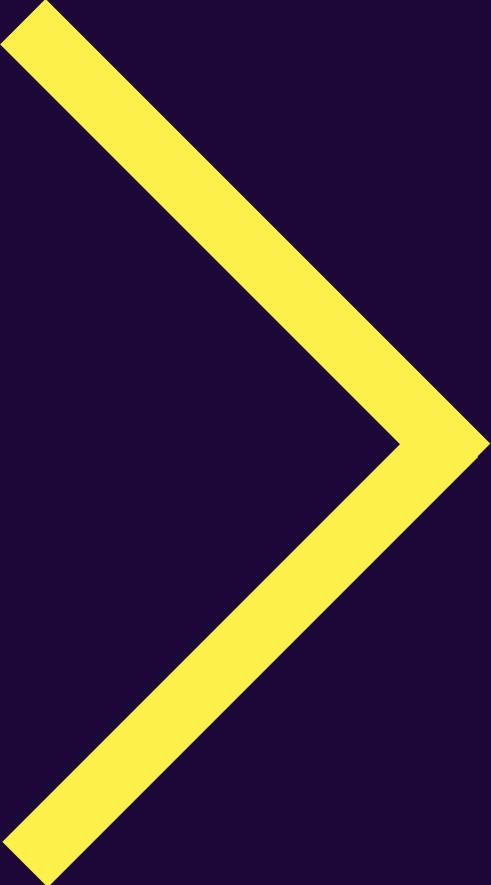


CI/CD Github Actions



buildah

Result



Models - Successful llms

Code intelligence advanced setup

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successful I fuzz tests
Phi 14b	100%	12m 8s	59.8k	0	0
Llama 3	100%	3m 41s	27.6k	0	0
Qwen 1.5 7b	89.47%	6m 9s	50k	0	0

Models - **Successful llms**

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 32b	43.08%	32m 57s	45.1k	2.04k	2
Gemma 3 27b	45.06%	33m 33s	40.2k	2.05k	2
Phi 14b	34.26%	36m 36s	71.5k	2.22k	1

Models - **Unsuccessful llms**

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
<i>Codellama 32b</i>	0.00%	50m 43s	74k	0	0
<i>Deepseek r1</i>	0.00%	1h 36m 53s	54.9k	0	0
<i>Deepseek code</i>	0.00%	1h 37m 39s	48.8k	0	0
<i>devstral</i>	0.00%	38m 57s	82.6k	0	0
<i>Llama 3 7b</i>	0.00%	27m 40s	268k	0	0
<i>Starcoder 2 15</i>	0.00%	19m 55s	114k	0	0
<i>Wizardcoder</i>	0.00%	32m 15s	32.5k	0	0
<i>Yi 34b</i>	0.00%	2h 39m 53s	74.9k	0	0
<i>Magistral 24b</i>	0.00%	-	-	0	0
<i>Mixtral</i>	0.00%	-	-	0	0

Models - *UnSuccessful llms*

fmt

<i>Models</i>	<i>Code Coverage</i>	<i>Time Taken</i>	<i>No of tokens used</i>	<i>Unique test cases</i>	<i>Successful I fuzz tests</i>
Qwen 2.5 coder 32b	0.00%	43m 20s	59.8k	0	0
Gemma 3 27b	0.00%	41m 59s	73.9k	0	0

Models - Successful llms

pugixml

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successful Ifuzz tests
Qwen 2.5 coder 32b	34.77%	37m 24s	43.1k	2.5k	1
Gemma 3 27b	0.00%	1h 12m 8s	122k	0	0

Models - Successful llms

Qwen 2.5 coder 32b

Repositories	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
jsoncons	45.64%	37m 24s	43.1k	2.5k	1
glm	0.00%	42m 55s	56.8k	0	0
spdlog	0.00%	49m 33s	62.9k	0	0

Finetuning

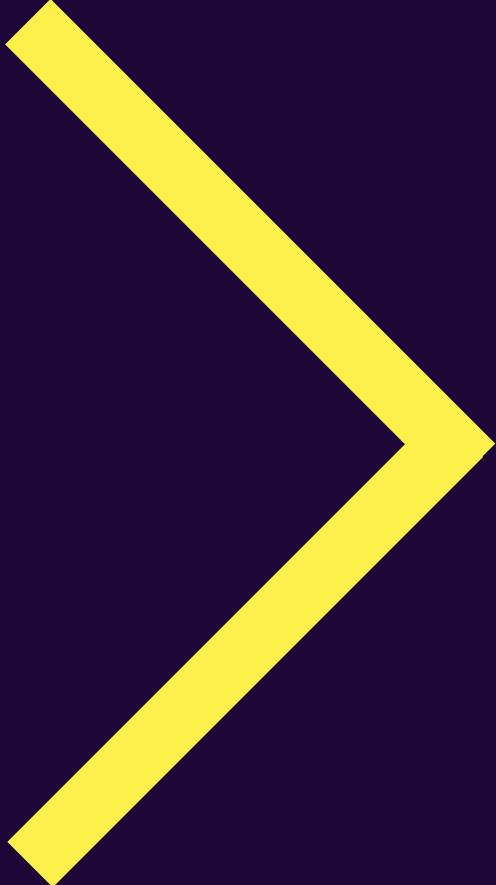
NAME	1	SIZE
• <i>qwen2.5-coder:1.5b</i>		<i>986 MB</i>
• <i>qwen-fuzzer-1.5-709-examples:latest</i>		<i>3.1 GB</i>
• <i>qwen-fuzzer-1.5-172-examples:latest</i>		<i>3.1 GB</i>

Finetuning - Successful llms

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successful I fuzz tests
Qwen 2.5 coder 1.5b	same	15 m	112k		
172 examples	same	12 m	65k		
709 examples	same	10 m	50k		

Costs



Literature review

Key Focus Areas:

- *AI techniques in software testing*
- *LLM-based fuzzing approaches*
- *CI/CD pipeline integration*
- *Automotive security applications*

From Traditional AI to LLMs - The Technical Journey

Key Milestones:

- **NeuFuzz (2019): First neural network-guided fuzzing - 35% improvement**
- **Deep RL Fuzzing (2018): 2.3× more crashes than AFL, 21 CVEs**
- **TitanFuzz (2022): Watershed moment - First LLM-based fuzzer**
- **Fuzz4All (2023): Universal multi-language fuzzing - 98 bugs found**

Performance Evolution:

- **2018-2020: Traditional ML approaches dominated (60%)**
- **2022-2025: LLM revolution with 50.84% coverage improvements**

The TitanFuzz Revolution & Beyond

TitanFuzz Impact (2022):

- ***Zero-shot capability without explicit constraints***
- ***30.38% higher coverage on TensorFlow***
- ***50.84% higher coverage on PyTorch***
- ***65 bugs discovered, 44 previously unknown***

Follow-up Advances:

- ***HGFuzzer: 24.8× speedup over traditional approaches***
- ***CKGFuzzer: Code knowledge graphs + LLMs***
- ***G²Fuzz: <\$0.2 for 24-hour fuzzing campaigns***

Automotive Domain Applications

Safety-Critical Breakthroughs:

- ***SAFLITE: Autonomous systems fuzzing - 234.5% improvement***
- ***CAN Bus AI Fuzzing: Real-time automotive network security***
- ***ECG Embedded OS: 32 new vulnerabilities in embedded systems***
- ***KernelGPT: 24 unknown bugs, 11 CVE assignments***

Automotive Applications:

- ***Neural network validation for autonomous driving***
- ***ECU firmware testing with real-time constraints***
- ***Multi-ECU system integration testing***

Research Gaps & Critical Findings

Identified Gaps:

- **Limited automotive-specific research - Most work targets general software**
- **Real-time constraint handling - Insufficient automotive timing requirements**
- **Multi-ECU testing - Lack of distributed architecture approaches**
- **Standardization - No safety-critical evaluation frameworks**

Quantitative Evidence:

- **20-90% consistent performance gains across all AI approaches**
- **2-24× speed improvements in specialized scenarios**
- **Hundreds of CVE discoveries with real-world impact**
- **Cost reduction: From \$1000s to <\$1 per campaign**

Key Takeaways

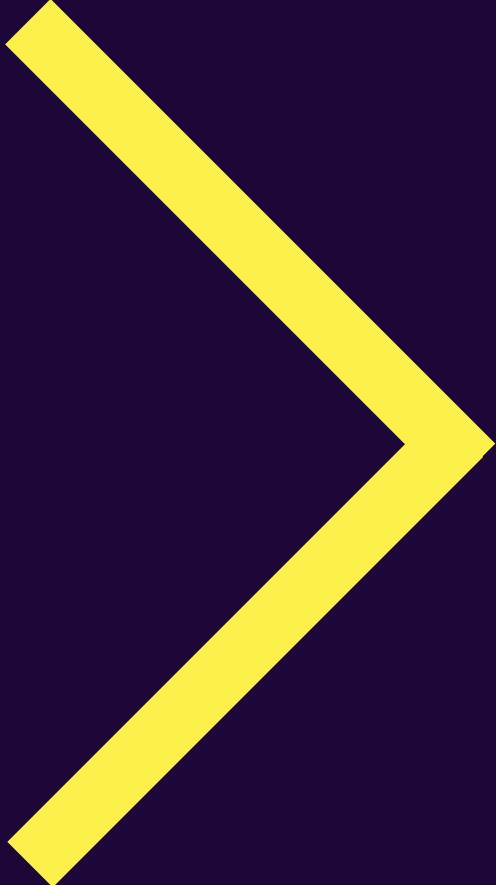
Major Findings:

- *AI-enhanced fuzzing consistently superior to traditional methods*
- *LLM-based approaches represent paradigm shift*
- *Industrial adoption accelerating with proven ROI*

Future Research Priorities:

- *Automotive-specific AI fuzzing frameworks*
- *Integration with safety standards (ISO 26262, ISO 21434)*
- *Real-time aware fuzzing techniques*

Costs



Costs

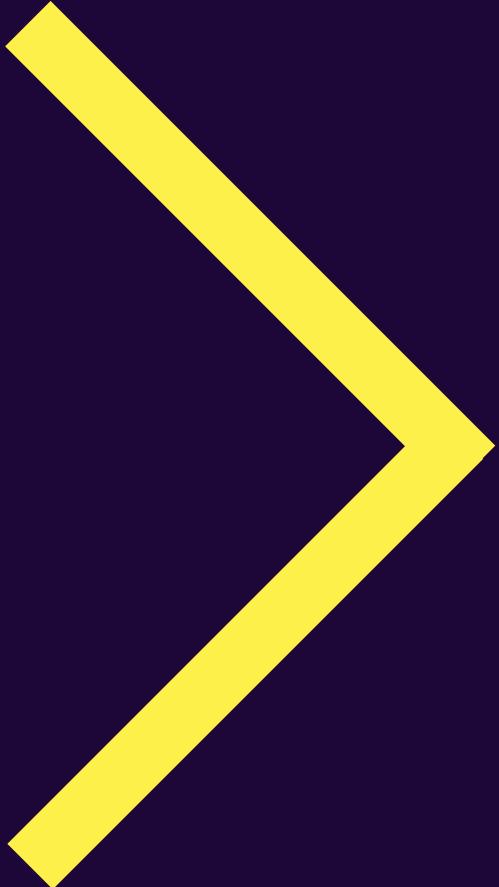
Daily Token Usage	Input Tokens	Output Tokens	Daily Cost	Monthly Cost (22 working days)	Annual Cost
Light Usage	5,000	7,500	€0.28	€6.16	€73.92
Moderate Usage	15,000	22,500	€0.83	€18.26	€219.12
Heavy Usage	50,000	75,000	€2.75	€60.50	€726.00
Enterprise Usage	100,000	150,000	€5.50	€121.00	€1,452.00

Costs

Test Scenario	Tokens Consumed	Cost per Run	Runs per Day	Daily Total
Single code file fuzzing	2,000 in + 3,000 out	€0.11	10	€1.10
Module testing	8,000 in + 12,000 out	€0.44	5	€2.20
Full application scan	25,000 in + 35,000 out	€1.30	2	€2.60
CI/CD pipeline integration	15,000 in + 20,000 out	€0.75	8	€6.00

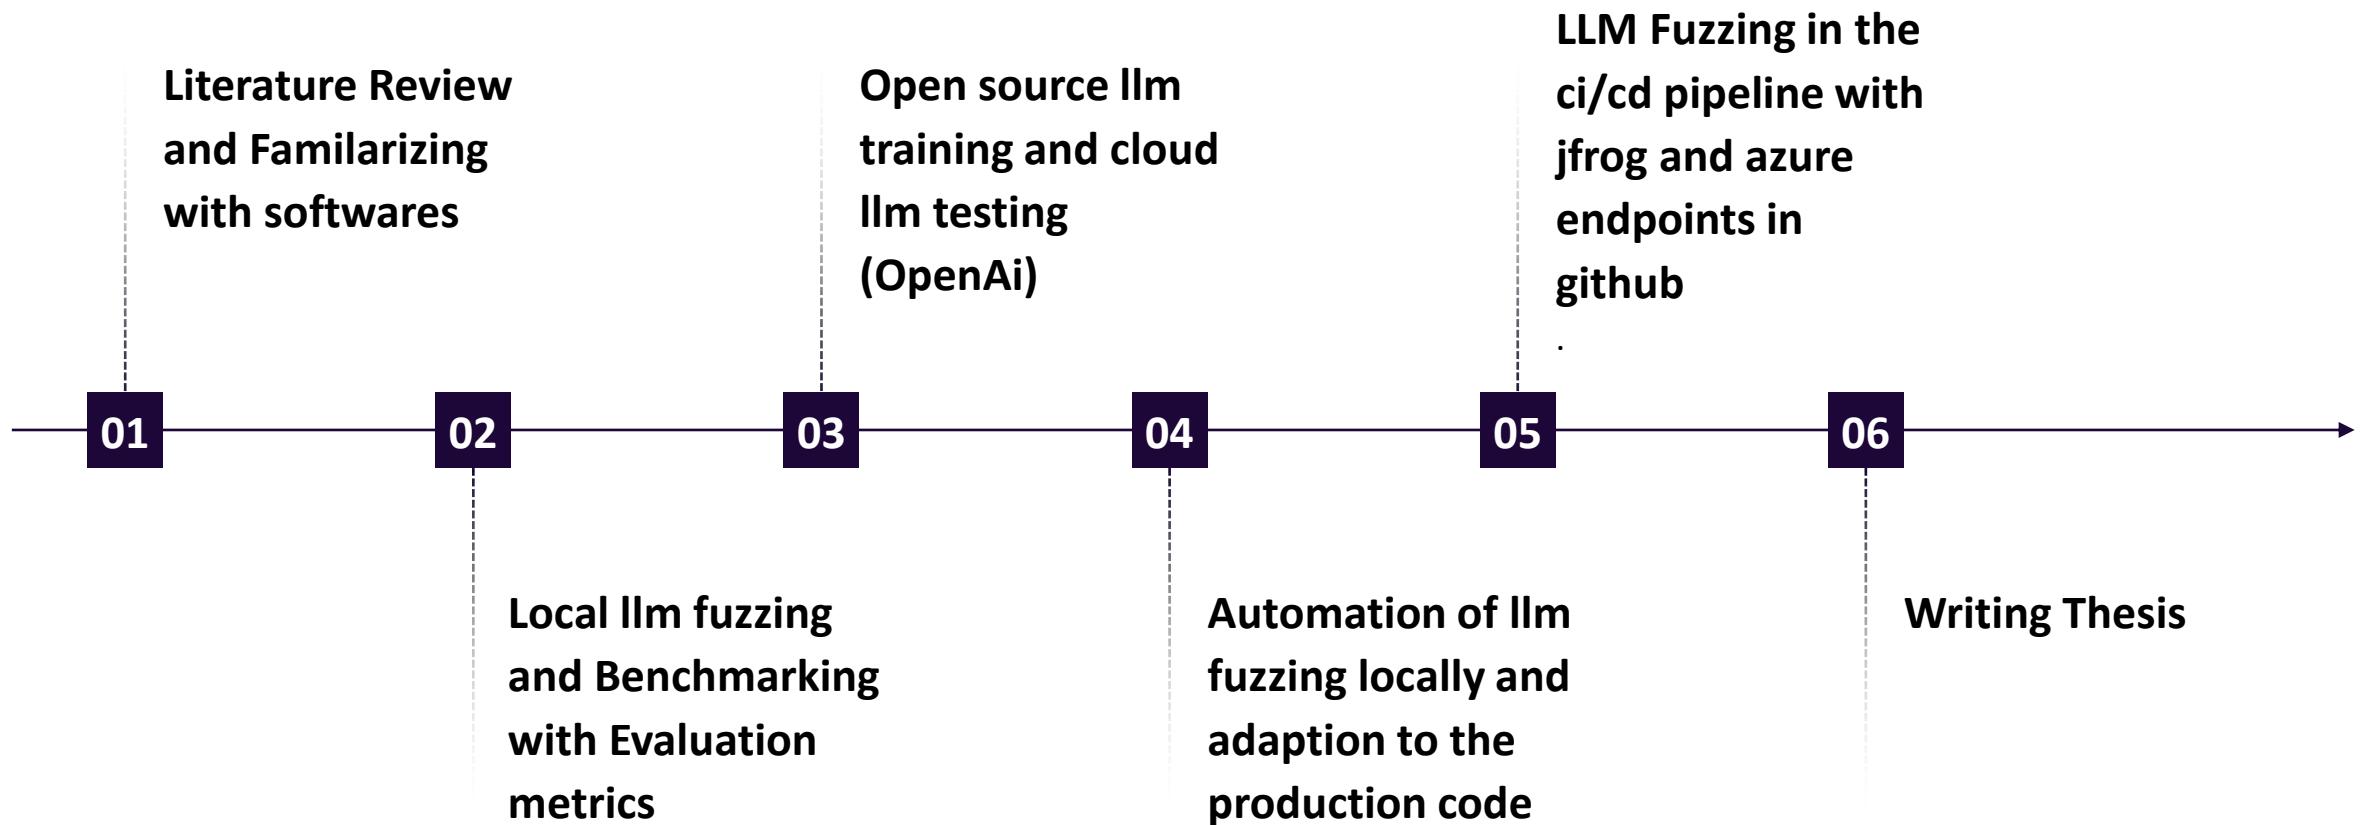
Rapidjson example

Conclusion



*This thesis focuses on
integrating AI and LLMs
into CI/CD/CT pipelines to
improve the security
testing of automotive
software*

Timeline



*Vielen Dank
für Ihre Aufmerksamkeit!*



Any Questions

