

CARIAD

Master Thesis: AI Usage in CI/CD/CT Pipelines for Compute Platforms in Automotives



We transform automotive mobility

CARIAD
A VOLKSWAGEN GROUP COMPANY

Agenda

// Previous Developments

// Status

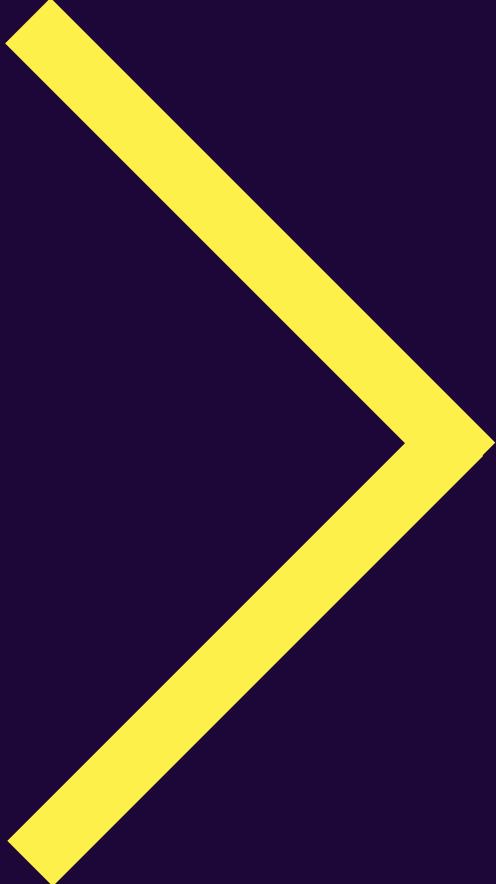
// Method

// Results

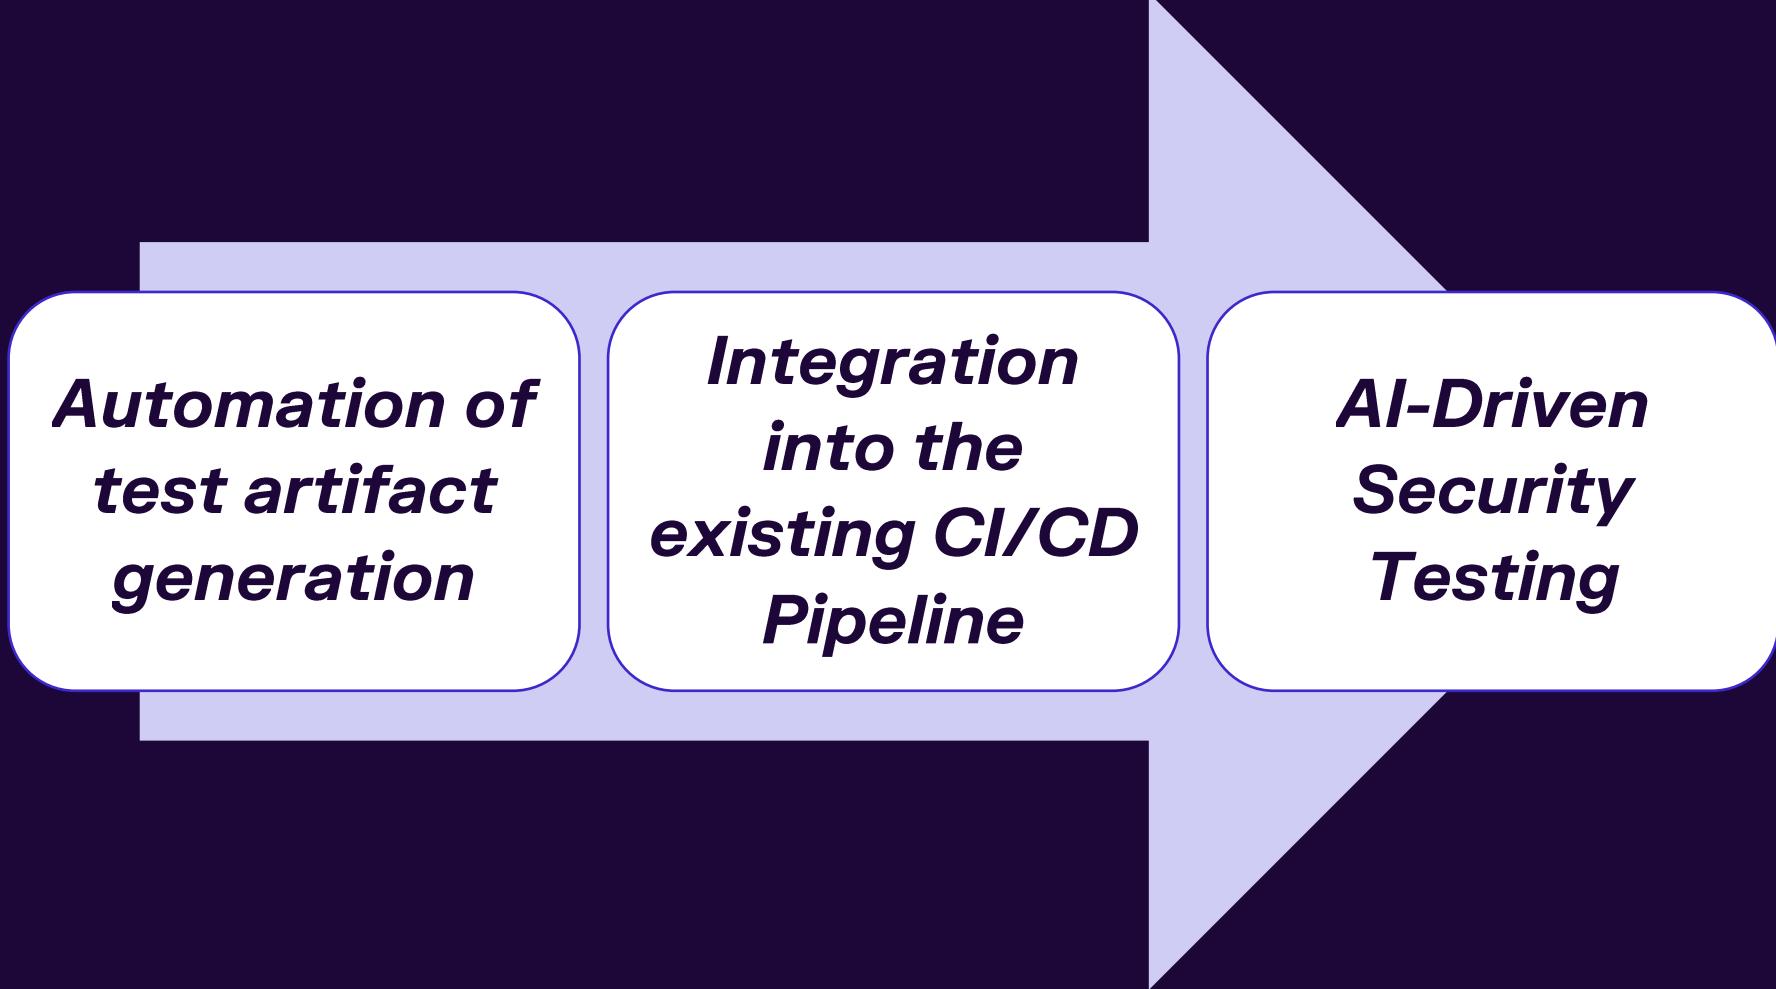
// Literature Review

// Next Steps

Previous Developments



Expected Outcomes



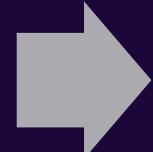
*Automation of
test artifact
generation*

*Integration
into the
existing CI/CD
Pipeline*

*AI-Driven
Security
Testing*

Previously

***Evaluation
Framework***



***Comparison
between local
llms***



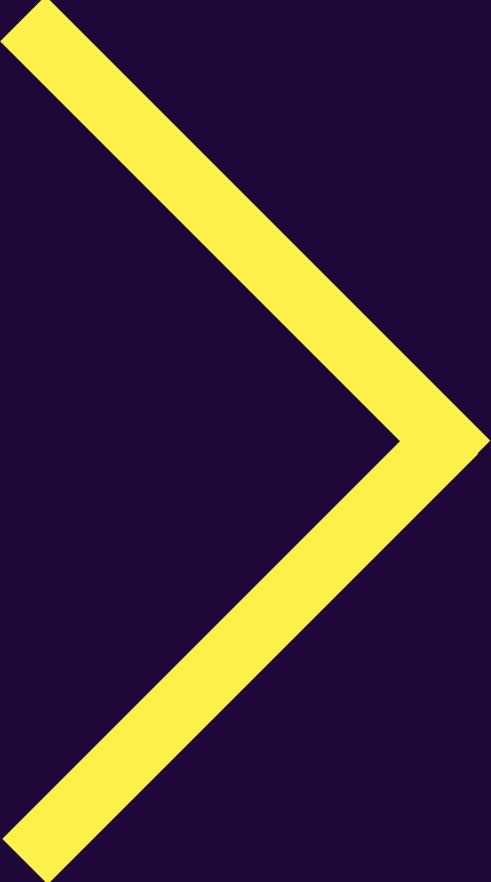
***Comparison
between local
and cloud llms***

Evaluation Framework

Metrics:

- ***Number of bugs found***
- ***Code coverage percentage***
- ***Time to generate fuzz tests***
- ***Number of tokens***
- ***Quality of generated code (measured via static analysis or manual review)***

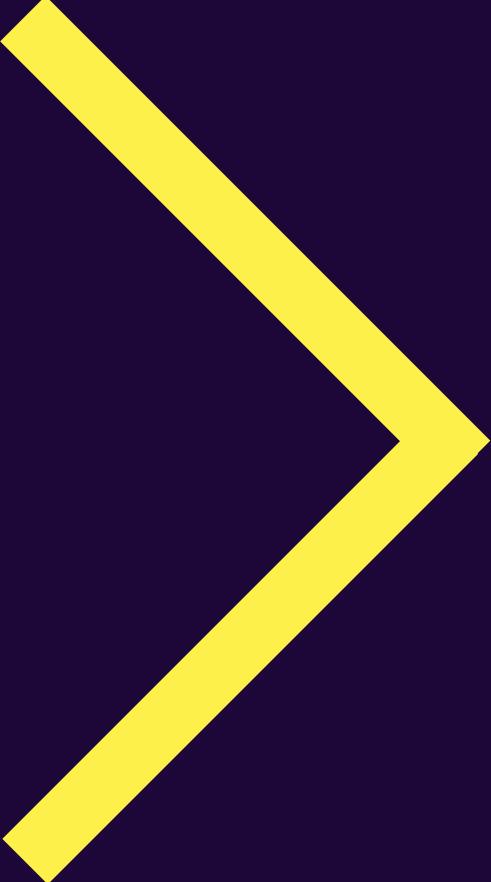
Status



Status

- *Different local setups*
- *Local llms comparison (14 llms)*
- *Tested on various cpp repositories and libraries (25)*
- *Automated workflow – half success*

Method



Method

**LLMs broad
classification**

n

**Code
specialized**

**General
purpose**

Method

LLM'S models and performance

*Small < 15B
parameters*

*Medium 15 –
25B parameters*

*Large 25 -35B
parameters*

*Extra Large
>40B
parameters*

Models

Large (25-35B parameters)

- ***gemma3:27b (27.4B)***
- ***qwen3:32b (32.8B)***
- ***deepseek-r1:32b (32.8B)***
- ***qwen2.5-coder:32b (32.8B)***
- ***deepseek-coder:33b (33B)***
- ***wizardcoder:33b (33B)***
- ***codellama:34b-instruct (34B)***
- ***yi:34b (34B)***

Models

Extra Large (> 40B parameters)

- ***mixtral:latest (46.7B)***

Previously Iterations

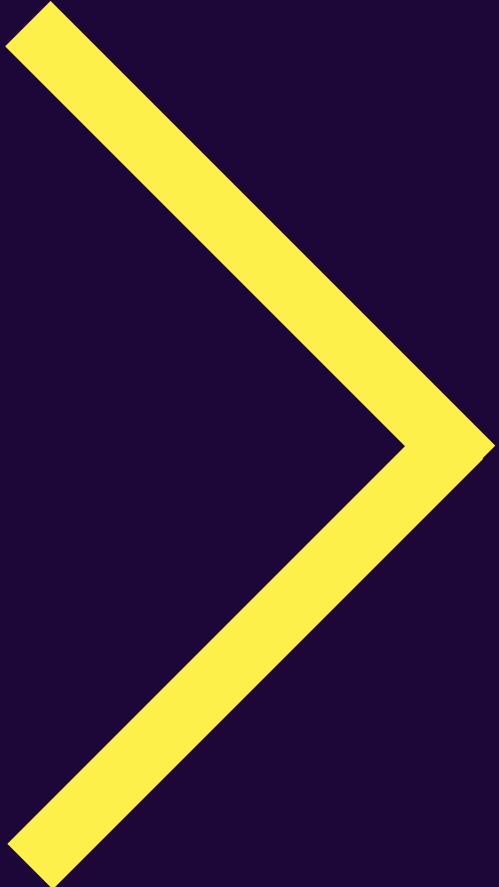
- ***qwen2.5-coder:7b (7B)***
- ***qwen2.5-coder:14b (14B)***
- ***qwen1.5:7b (7B)***

Repositories



repositories

Results



Models - Successful llms

Code intelligence advanced setup

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Phi 14b	100%	12m 8s	59.8k	0	0
Llama 3	100%	3m 41s	27.6k	0	0
Qwen 1.5 7b	89.47%	6m 9s	50k	0	0

Models - *Successful llms*

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
<i>Qwen 2.5 coder 32b</i>					
<i>Gemma 3 27b</i>	45.06%	33m 33s	40.2k	2.05k	2
<i>Phi 14b</i>	34.26%	36m 36s	71.5k	2.22k	1

Models - *Unsuccessful llms*

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
<i>Codellama 32b</i>	0.00%	50m 43s	74k	0	0
<i>Deepseek r1</i>	0.00%	1h 36m 53s	54.9k	0	0
<i>Deepseek code</i>	0.00%	1h 37m 39s	48.8k	0	0
<i>devstral</i>	0.00%	38m 57s	82.6k	0	0
<i>Llama 3 7b</i>	0.00%	27m 40s	268k	0	0
<i>Starcoder 2 15</i>	0.00%	19m 55s	114k	0	0
<i>Wizardcoder</i>	0.00%	32m 15s	32.5k	0	0
<i>Yi 34b</i>	0.00%	2h 39m 53s	74.9k	0	0
<i>Magistral 24b</i>	0.00%	-	-	0	0
<i>Mixtral</i>	0.00%	-	-	0	0

Models - *UnSuccessful llms*

fmt

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 32b	0.00%	43m 20s	59.8k	0	0
Gemma 3 27b	0.00%	41m 59s	73.9k	0	0

Models - *Successful llms*

pugixml

<i>Models</i>	<i>Code Coverage</i>	<i>Time Taken</i>	<i>No of tokens used</i>	<i>Unique test cases</i>	<i>Successfull fuzz tests</i>
Qwen 2.5 coder 32b	34.77%	37m 24s	43.1k	2.5k	1
Gemma 3 27b	0.00%	1h 12m 8s	122k	0	0

Models - *Successful llms*

Qwen 2.5 coder 32b

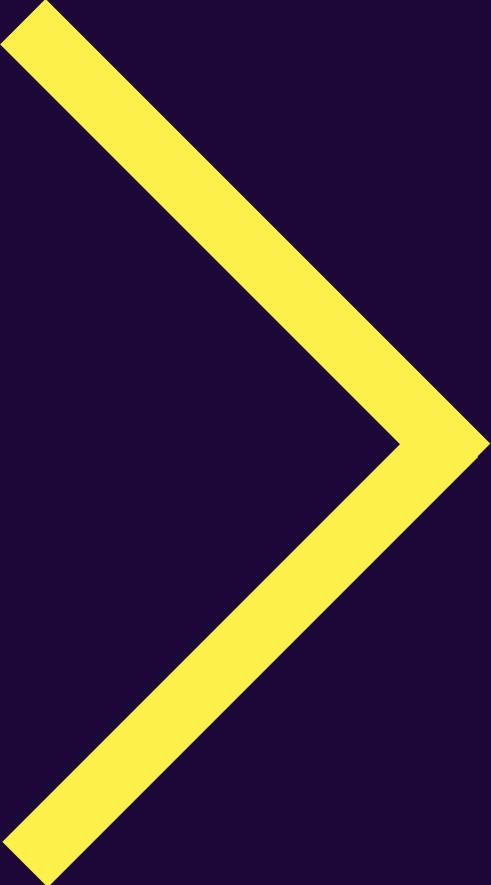
<i>Repositories</i>	<i>Code Coverage</i>	<i>Time Taken</i>	<i>No of tokens used</i>	<i>Unique test cases</i>	<i>Successfull fuzz tests</i>
jsoncons	45.64%	37m 24s	43.1k	2.5k	1
glm	0.00%	42m 55s	56.8k	0	0
spdlog	0.00%	49m 33s	62.9k	0	0

Literature Review

- 1. How effective are they? Exploring large language model based fuzz driver generation**
- 2. On the Challenges of Fuzzing Techniques via Large Language Models**

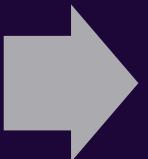
Literature Review

Next Steps



Next Steps

***Training LLM's
with code
examples***



***Understanding
cloud LLM's and
finetuning with
code examples***



***Test Artifacts
Generation***

Finetuning local llms

Qwen 2.5 coder 32b instruct full model 60gb

LoRA (Low-Rank Adaptation)

- *LoRA rank (r): 16 – controls adaptation size*
- *LoRA alpha: 32 – scaling for adaptation*
- *Dropout: 0.1 – prevents overfitting*
- *Target modules: q_proj, v_proj, etc. – efficient fine-tuning*
- *Device: auto – runs on best available hardware*
- *Dtype: float16 – faster, less memory*
- *Efficient model loading and saving (safetensors)*

Code examples for training

Jsonl format

```
{"prompt": "Generate a comprehensive parser fuzzer for structured data\n\nRequirements:\n- Maximize code coverage and edge case detection\n- Generate intelligent, targeted inputs\n- Optimize for vulnerability discovery speed\n- Include comprehensive error handling."}
```

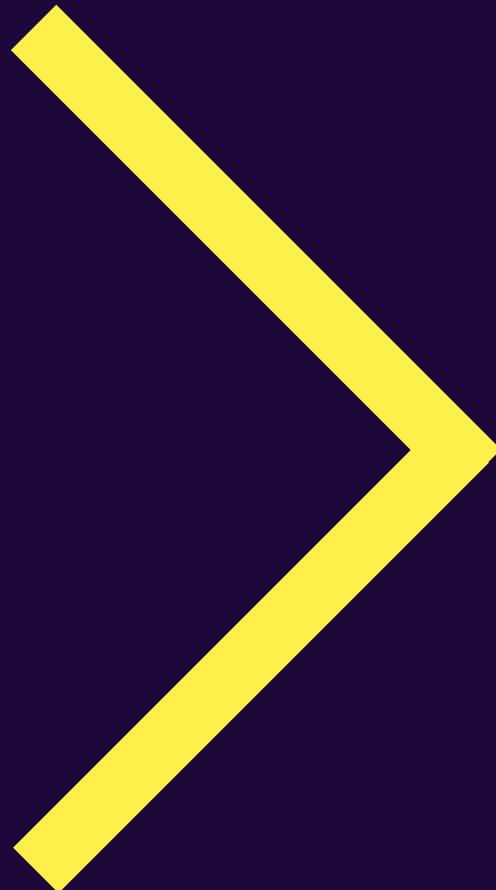
Code examples for training

Qwen 2.5 coder 32b

"completion":

```
void train_timing_attack(const uint8_t* data, size_t size) {\n    if (size < 20)\n        return;\n\n    const char* secret = "supersecretpassword123";\n    std::string user_input(reinterpret_cast<const char*>(data),\n                           std::min(size,\n                                   strlen(secret)));\n\n    auto start =\n        std::chrono::high_resolution_clock::now();\n\n    // VULNERABLE: Early return creates timing difference\n    bool matches = true;\n    for (size_t i = 0; i < strlen(secret) && i < user_input.length(); i++) {\n        if (secret[i] != user_input[i]) {\n            matches = false;\n            break; // Early return reveals timing information\n        }\n        // Simulate processing time\n    }\n\n    auto end =\n        std::chrono::high_resolution_clock::now();\n    auto duration =\n        std::chrono::duration_cast<std::chrono::microseconds>(end -\n                                                               start);\n\n    printf("Password check took %ld microseconds, match: %s\\n\\n",\n          duration.count(), matches ? "true" : "false");\n}
```

Thank you!



Any Questions

