

CARIAD

Master Thesis: AI Usage in CI/CD/CT Pipelines for Compute Platforms in Automotives



We transform automotive mobility

CARIAD
A VOLKSWAGEN GROUP COMPANY

Agenda

// Introduction

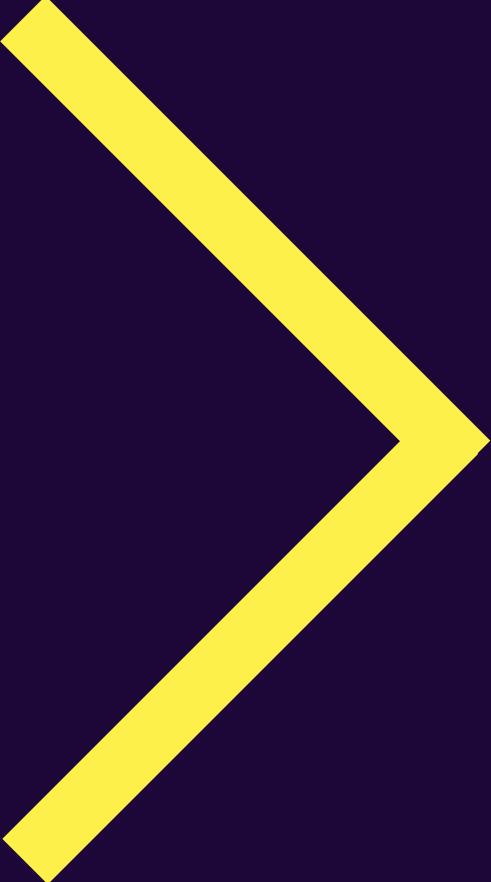
// Method

// Result

/7 Status

// Next Steps

Introduction



Introduction

- ***Rapidly growing software complexity & shorter release cycles***
- ***Automotive ECUs are safety-critical, hence zero tolerance***
- ***Traditional security tests cannot keep pace with CI/CD demand***

Problem Statement

- ***Current white-box fuzzing & testing are manual or slow to scale***
- ***Vulnerabilities may slip through nightly CI due to time limits***
- ***Need an AI-guided approach integrated into CI/CD/CT to***
 - ***boost path coverage***
 - ***reduce manual fuzz test case creation***
 - ***auto-generate actionable test artifacts***

Research Objectives

Technique Design

- *AI-assisted white-box fuzzing for automotive targets*

Pipeline Integration

- *Embed continuous fuzzing into existing CI/CD/CT*

Artifact & Impact Automation

- *Auto-generate test cases, reports, quality matrix*
- *Measure coverage, MTTV, and CI latency vs. baseline*

Expected Outcomes

*AI Generated
fuzz test
cases and
intelligent
mutation*

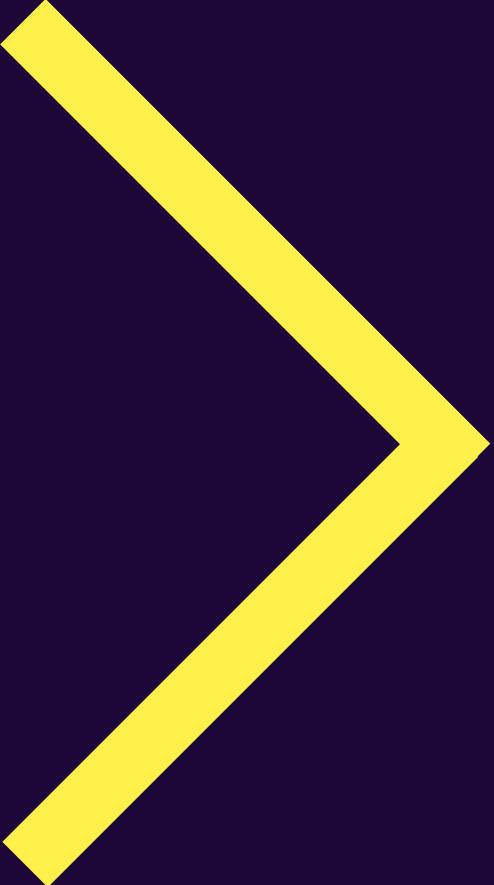
*Automation of
test artifact
generation*

*Integration
into the
existing CI/CD
pipeline*

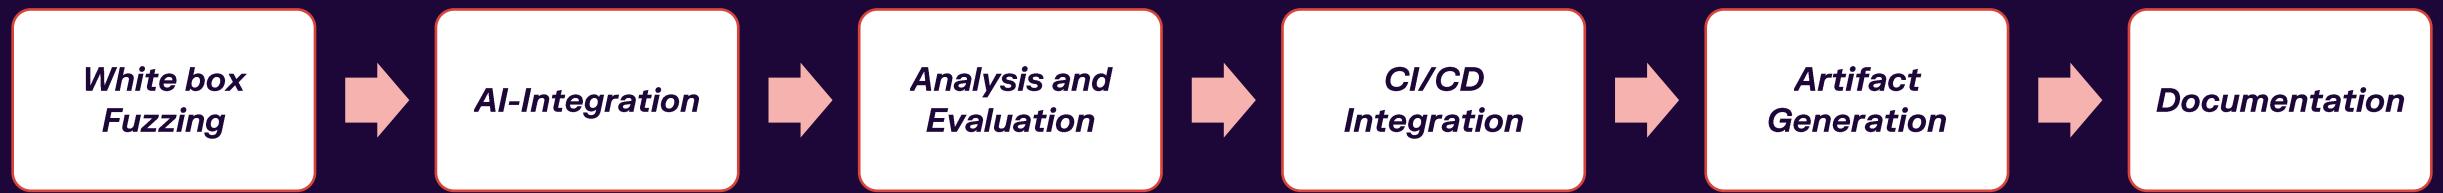
*AI-Driven
Security
Testing*

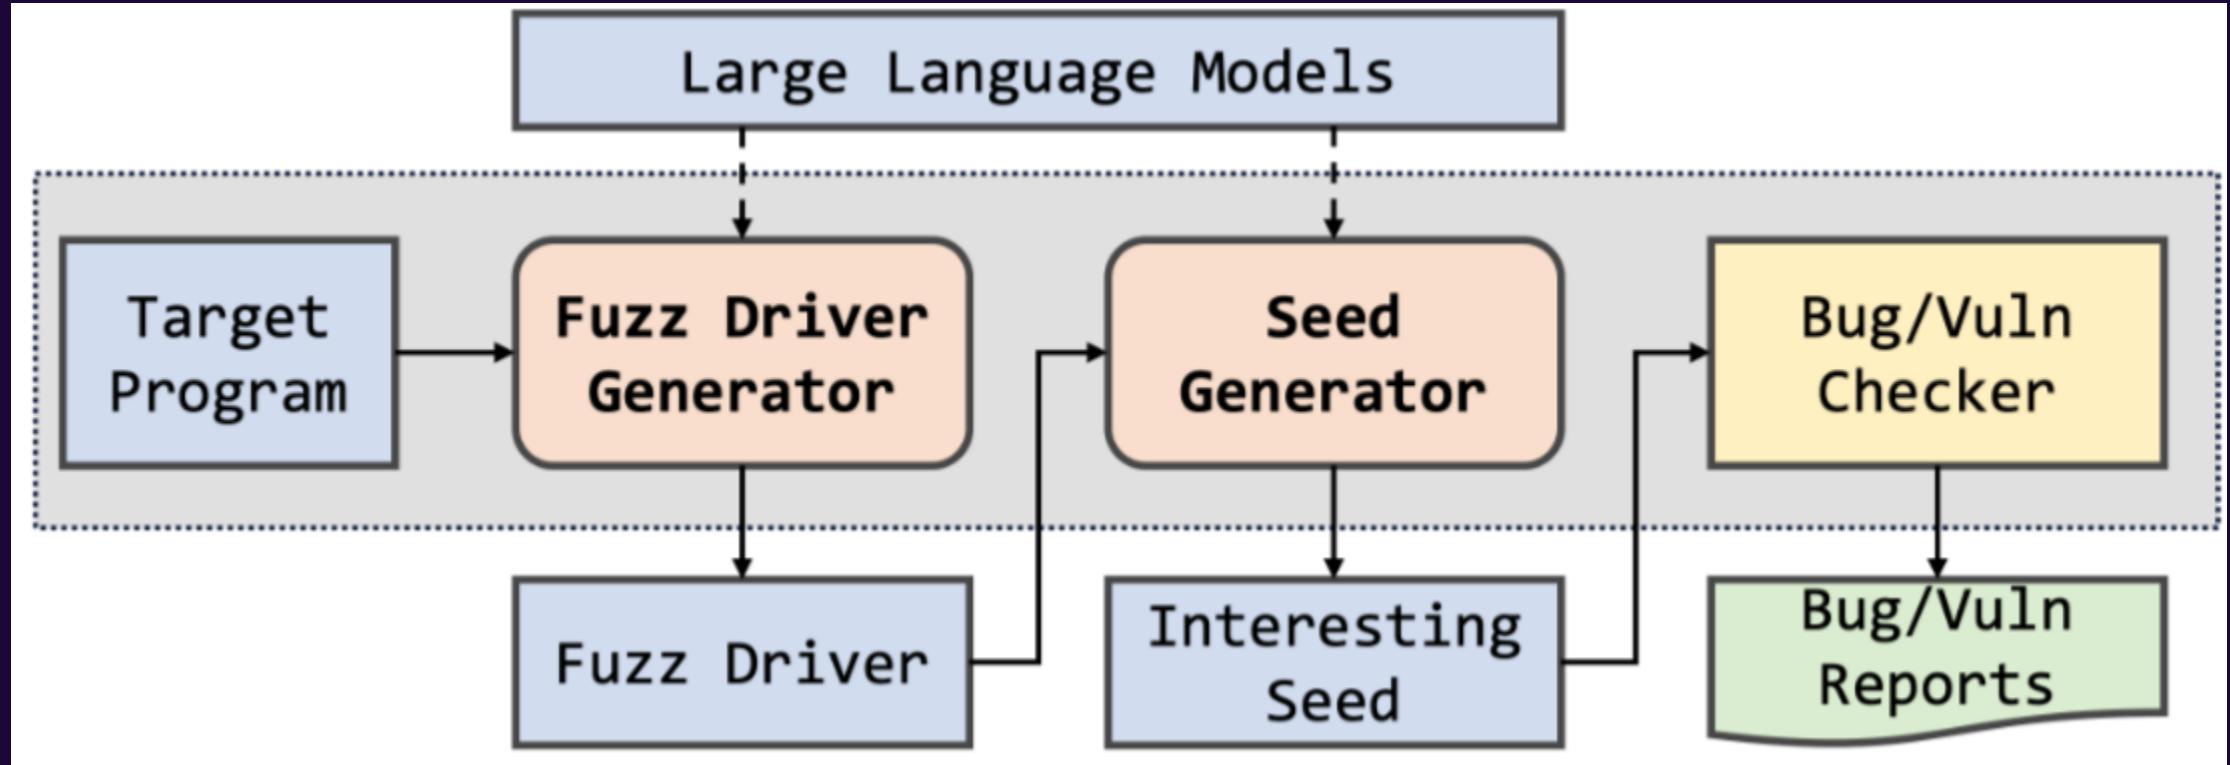
May

Method



Overview





Method

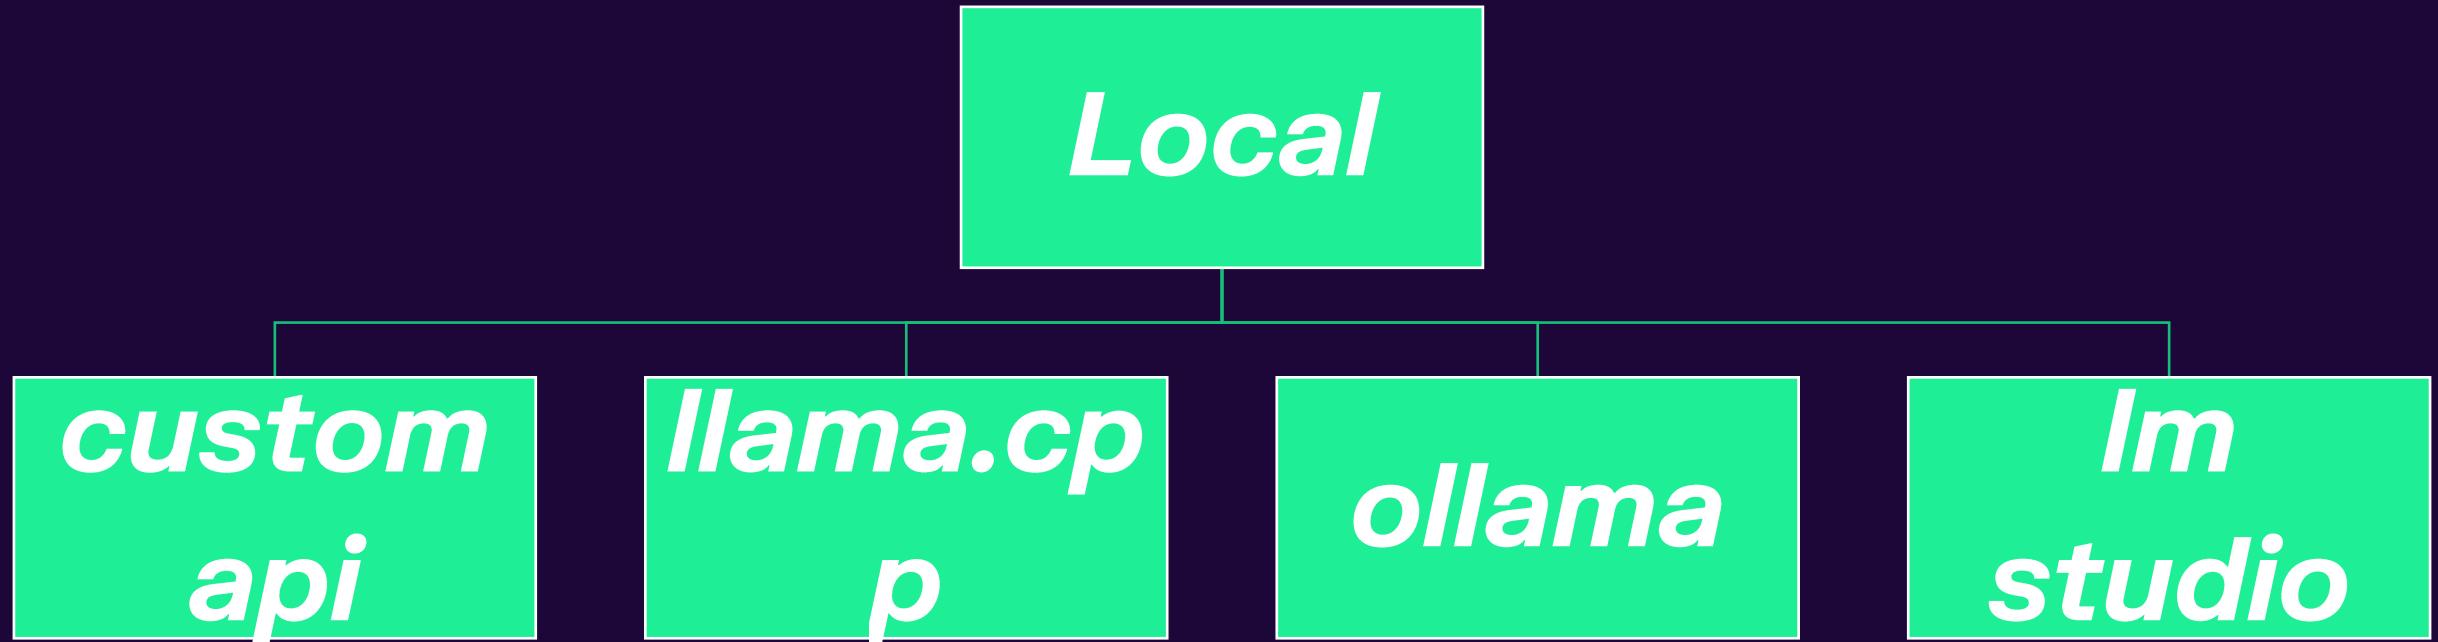
LLMs

The diagram illustrates the relationship between Large Language Models (LLMs) and their deployment methods. A central box labeled "LLMs" branches down into two separate boxes: "local" on the left and "api" on the right, representing the two primary ways to utilize these models.

local

api

Method



Method

*broad
classification*

*Code
specialized*

*General
purpose*

Method

LLM'S *models and performance*

***Small < 15B
parameters***

***Medium 15 -
25B
parameters***

***Large 25 -
35B
parameters***

***Extra Large
>40B
parameters***

Models

Large (25-35B parameters)

- ***gemma3:27b (27.4B)***
- ***qwen3:32b (32.8B)***
- ***deepseek-r1:32b (32.8B)***
- ***qwen2.5-coder:32b (32.8B)***
- ***deepseek-coder:33b (33B)***
- ***wizardcoder:33b (33B)***
- ***codellama:34b-instruct (34B)***
- ***yi:34b (34B)***

Models

Extra Large (> 40B parameters)

- ***mixtral:latest (46.7B)***

Previously Iterations

- ***qwen2.5-coder:7b (7B)***
- ***qwen2.5-coder:14b (14B)***
- ***qwen1.5:7b (7B)***

Repositories



repositories

Evaluation Framework

Metrics:

- *Number of bugs found*
- *Code coverage percentage*
- *Time to generate fuzz tests*
- *Quality of generated code (measured via static analysis or manual review)*
- *Number of tokens used*

Finetuning local llms

Qwen 2.5 coder 32b instruct full model 60gb

LoRA (Low-Rank Adaptation)

- *LoRA rank (r): 16 – controls adaptation size*
- *LoRA alpha: 32 – scaling for adaptation*
- *Dropout: 0.1 – prevents overfitting*
- *Target modules: q_proj, v_proj, etc. – efficient fine-tuning*
- *Device: auto – runs on best available hardware*
- *Dtype: float16 – faster, less memory*
- *Efficient model loading and saving (safetensors)*

Code examples for training

Jsonl format

```
{"prompt": "Generate a comprehensive parser fuzzer for structured data\n\nRequirements:\n- Maximize code coverage and edge case detection\n- Generate intelligent, targeted inputs\n- Optimize for vulnerability discovery speed\n- Include comprehensive error handling."}
```

Code examples for training

Qwen 2.5 coder 32b

"completion":

```
"void train_timing_attack(const uint8_t* data, size_t size) {\n    if (size < 20) return;\n\n    const char* secret = \"supersecretpassword123\";\n    std::string user_input(reinterpret_cast<const char*>(data),\n                           std::min(size, strlen(secret)));\n\n    auto start =\n        std::chrono::high_resolution_clock::now();\n\n    // VULNERABLE: Early return creates timing difference\n    bool matches = true;\n    for (size_t i = 0; i < strlen(secret) && i < user_input.length(); i++) {\n        if (secret[i] != user_input[i]) {\n            matches = false;\n            break; // Early return reveals timing information\n        }\n        // Simulate processing time\n    }\n\n    auto end =\n        std::chrono::high_resolution_clock::now();\n    auto duration =\n        std::chrono::duration_cast<std::chrono::microseconds>(end - start);\n\n    printf(\"Password check took %ld microseconds, match: %s\\n\", duration.count(), matches ? \"true\" : \"false\");\n}
```

Finetuning

- ***32b – needs more than 60gb ram***
- ***14b – (32 – 37gb) ram***
- ***7b – 24gb ram without any background process***
- ***1.5b – 14gb ram***

Finetuning

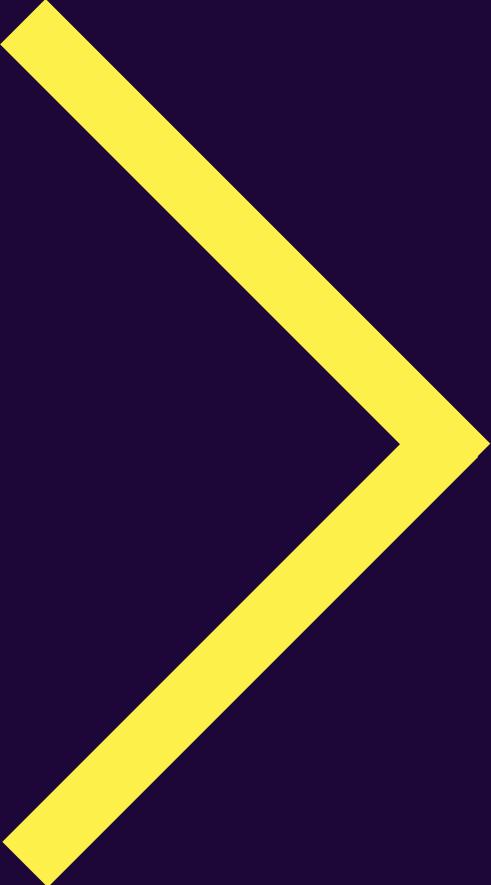
<i>NAME</i>	<i>SIZE</i>
• <i>qwen2.5-coder:1.5b</i>	<i>986 MB</i>
• <i>qwen-fuzzer-1.5-709-examples:latest</i>	<i>3.1 GB</i>
• <i>qwen-fuzzer-1.5-172-examples:latest</i>	<i>3.1 GB</i>

Finetuning - Successful llms

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 1.5b	same	15 m	112k		
172 examples	same	12 m	65k		
709 examples	same	10 m	50k		

Result



Models - Successful llms

Code intelligence advanced setup

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Phi 14b	100%	12m 8s	59.8k	0	0
Llama 3	100%	3m 41s	27.6k	0	0
Qwen 1.5 7b	89.47%	6m 9s	50k	0	0

Models - *Successful llms*

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 32b	43.08%	32m 57s	45.1k	2.04k	2
Gemma 3 27b	45.06%	33m 33s	40.2k	2.05k	2
Phi 14b	34.26%	36m 36s	71.5k	2.22k	1

Models - *Unsuccessful llms*

Yaml cpp (35 files, 1061 candidates)

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
<i>Codellama 32b</i>	0.00%	50m 43s	74k	0	0
<i>Deepseek r1</i>	0.00%	1h 36m 53s	54.9k	0	0
<i>Deepseek code</i>	0.00%	1h 37m 39s	48.8k	0	0
<i>devstral</i>	0.00%	38m 57s	82.6k	0	0
<i>Llama 3 7b</i>	0.00%	27m 40s	268k	0	0
<i>Starcoder 2 15</i>	0.00%	19m 55s	114k	0	0
<i>Wizardcoder</i>	0.00%	32m 15s	32.5k	0	0
<i>Yi 34b</i>	0.00%	2h 39m 53s	74.9k	0	0
<i>Magistral 24b</i>	0.00%	-	-	0	0
<i>Mixtral</i>	0.00%	-	-	0	0

Models - *UnSuccessful llms*

fmt

Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 32b	0.00%	43m 20s	59.8k	0	0
Gemma 3 27b	0.00%	41m 59s	73.9k	0	0

Models - *Successful llms*

pugixml

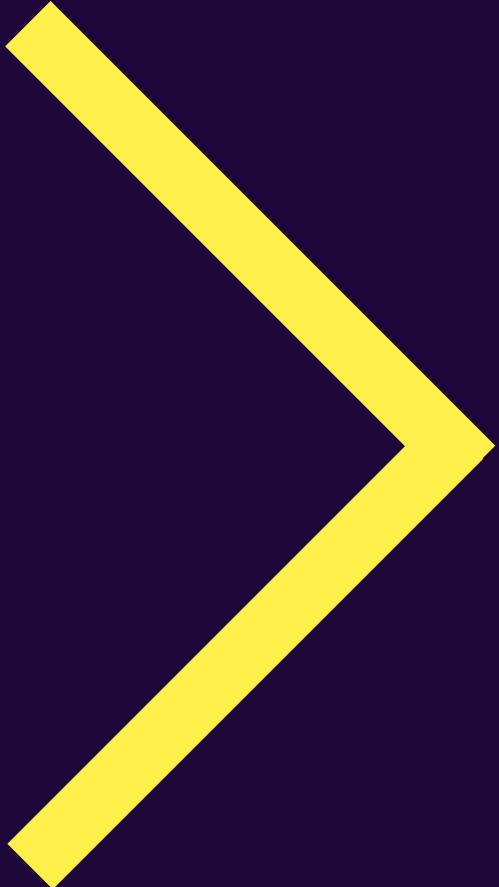
Models	Code Coverage	Time Taken	No of tokens used	Unique test cases	Successfull fuzz tests
Qwen 2.5 coder 32b	34.77%	37m 24s	43.1k	2.5k	1
Gemma 3 27b	0.00%	1h 12m 8s	122k	0	0

Models - *Successful llms*

Qwen 2.5 coder 32b

<i>Repositories</i>	<i>Code Coverage</i>	<i>Time Taken</i>	<i>No of tokens used</i>	<i>Unique test cases</i>	<i>Successfull fuzz tests</i>
jsoncons	45.64%	37m 24s	43.1k	2.5k	1
glm	0.00%	42m 55s	56.8k	0	0
spdlog	0.00%	49m 33s	62.9k	0	0

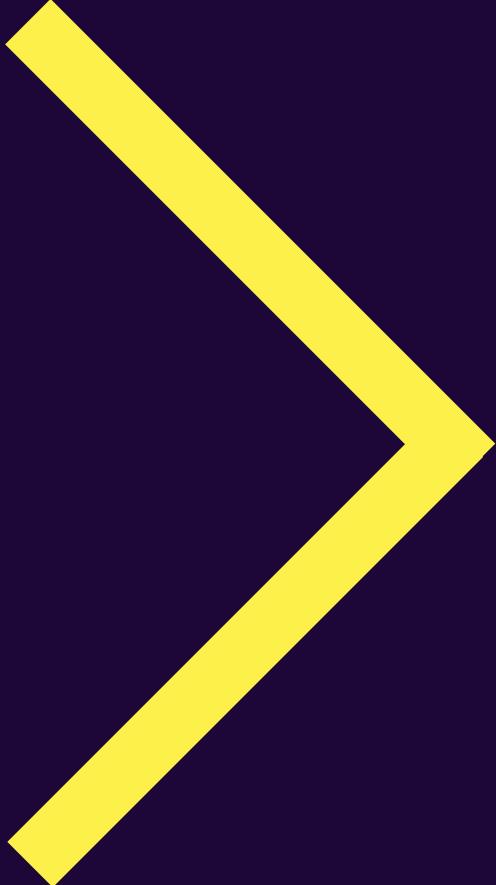
Status



Status

- *CI/CD Integration in github actions without spark*
- *Automation of spark with open source libraries locally*

Final Steps



Next steps

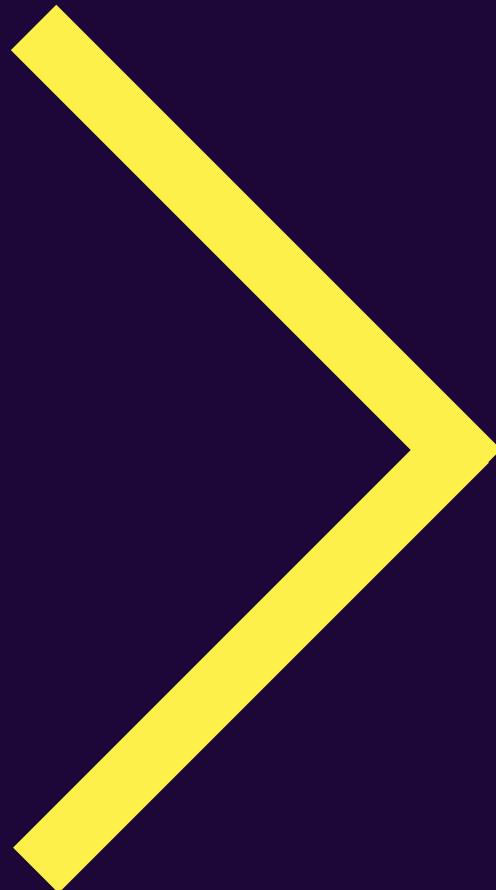
*Ci spark integration
in github actions*

*Adapting the ci
spark for company
adapted (e3 sdk
open source)*

Documentation

*This thesis focuses on integrating AI
and LLMs into CI/CD/CT pipelines to
improve the security testing of
automotive software*

Thank you!



Any Questions

