

Master Thesis: AI Usage in CI/CD/CT Pipelines for Compute Platforms in Automotive


Morris Darren Babu

Friedrich-Alexander-Universität Erlangen-Nürnberg, Hardware-Software-Co-Design

Sept 22, 2025


Agenda

- 01 Introduction & Problem Statement
- 02 Research Objectives
- 03 Research Questions
- 04 Literature Review & State of the Art
- 05 Overview
- 06 Method
- 07 conclusion



Introduction & Problem Statement


- Rapidly growing software complexity & shorter release cycles
- Automotive ECUs are safety-critical, hence zero tolerance
- Traditional security tests cannot keep pace with CI/CD demand



The Challenges



**100M→300M
lines by 2030**



**530
vulnerabilities
in 2024**



**97%
remote attacks**


- Current white-box fuzzing & testing are manual or slow to scale
- Vulnerabilities slip through nightly CI builds due to time constraints
- Manual fuzz driver creation requires deep expertise - bottleneck for development teams
- Growing attack surface with connected and autonomous vehicle features

Research Objectives

- **AI-assisted white-box fuzzing** for automotive compute platforms
- **Novel security testing methods** integrating LLMs with fuzzing capabilities
- **Automated test artifact generation** - test cases, procedures, reports, quality matrices

Pipeline Integration

- **Seamless CI/CD/CT integration** with existing automotive development workflows
- **Continuous fuzzing** embedded into daily development cycles
- **Zero-downtime deployment** of AI-enhanced security testing



- **Measure coverage, MTTV, and CI latency** vs. baseline traditional methods
- **Comprehensive cost-benefit analysis** for enterprise adoption
- **Real-world validation** within CARIAD's development environment

AI Generated
fuzz test
cases and
intelligent
mutation

Automation
of test artifact
generation

Integration
into the
existing
CI/CD
pipeline

AI-Driven
Security
Testing

Research Questions

Primary Research Questions

RQ1: Effectiveness of LLM-Generated Fuzz Drivers

Can Large Language Models generate effective fuzz drivers that achieve code coverage and vulnerability detection performance comparable to manually-written drivers?

RQ2: Model Selection and Optimization

Which LLM architectures, training approaches, and optimization techniques are most effective for automotive fuzzing applications?

RQ3: CI/CD Integration Feasibility

Can AI-driven fuzzing be integrated into automotive CI/CD/CT pipelines without significantly impacting development velocity or resource utilization?

Secondary Research Questions

RQ4: Automated Test Artifact Generation

Can LLMs automatically generate comprehensive test artifacts (test procedures, reports, quality matrices) that meet automotive development and compliance requirements?

RQ5: Continuous Learning and Adaptation

How can AI-driven fuzzing systems continuously improve their effectiveness through learning from fuzzing campaigns, bug discoveries, and developer feedback?

RQ6: Cost-Effectiveness and ROI

What are the economic implications of AI-enhanced fuzzing deployment, and under what conditions does it provide positive return on investment for automotive organizations?

Literature Review & State of the Art

Key Focus Areas:

- AI techniques in software testing
- LLM-based fuzzing approaches
- CI/CD pipeline integration
- Automotive security applications

Key Milestones:

- NeuFuzz (2019): First neural network-guided fuzzing - **35% improvement**
- Deep RL Fuzzing (2018): **2.3× more crashes than AFL, 21 CVEs**
- TitanFuzz (2022): **Watershed moment** - First LLM-based fuzzer
- Fuzz4All (2023): Universal multi-language fuzzing - **98 bugs found**

Performance Evolution:

- **2018-2020:** Traditional ML approaches dominated (60%)
- **2022-2025:** LLM revolution with **50.84% coverage improvements**

TitanFuzz Impact (2022):

- **Zero-shot capability** without explicit constraints
- **30.38% higher coverage** on TensorFlow
- **50.84% higher coverage** on PyTorch
- **65 bugs discovered, 44 previously unknown**

Follow-up Advances:

- HGFuzzer: **24.8× speedup** over traditional approaches
- CKGFuzzer: Code knowledge graphs + LLMs
- G²Fuzz: **<\$0.2 for 24-hour fuzzing** campaigns

Safety-Critical Breakthroughs:

- SAFLITE: Autonomous systems fuzzing - **234.5% improvement**
- CAN Bus AI Fuzzing: Real-time automotive network security
- ECG Embedded OS: **32 new vulnerabilities** in embedded systems
- KernelGPT: **24 unknown bugs, 11 CVE assignments**

Automotive Applications:

- Neural network validation for autonomous driving
- ECU firmware testing with real-time constraints
- Multi-ECU system integration testing

Identified Gaps:

- **Limited automotive-specific research** - Most work targets general software
- **Real-time constraint handling** - Insufficient automotive timing requirements
- **Multi-ECU testing** - Lack of distributed architecture approaches
- **Standardization** - No safety-critical evaluation frameworks

Quantitative Evidence:

- **20-90% consistent performance gains** across all AI approaches
- **2-24× speed improvements** in specialized scenarios
- **Hundreds of CVE discoveries** with real-world impact
- **Cost reduction:** From \$1000s to <\$1 per campaign

Major Findings:


- **AI-enhanced fuzzing consistently superior** to traditional methods
- **LLM-based approaches represent paradigm shift**
- **Industrial adoption accelerating** with proven ROI


Future Research Priorities:

- Automotive-specific AI fuzzing frameworks
- Integration with safety standards (ISO 26262, ISO 21434)
- Real-time aware fuzzing techniques


Overview

Overview






Process Flow Diagram for Automotive AI Fuzzing




Tools used:

- Visual Studio Code
- CiFuzz (exe or docker image)
- Cmake Build systems and libfuzzer
- Azure Open AI endpoint
- Github Actions (CI/CD) and Jenkins
- Jfrog and Artifactory (docker images)




Method




LLMs broad classification

Code
specialized


General
purpose







Bugs found




Code coverage %



Fuzz test generation time



Generated code quality




Tokens used

LoRA (Low-Rank Adaptation)

- LoRA rank (r): 16 – controls adaptation size
- LoRA alpha: 32 – scaling for adaptation
- Dropout: 0.1 – prevents overfitting
- Target modules: q_proj, v_proj, etc. – efficient fine-tuning
- Device: auto – runs on best available hardware
- Dtype: float16 – faster, less memory
- Efficient model loading and saving (safetensors)

LoRA for LLM Fine-Tuning



x : input

Frozen W : frozen weights

A, B: low-rank matrices

Scale arrow: scaling factor

Limitations

To train open source LLMS requirements

- 32b – needs more than 60gb ram
- 14b – (32 – 37gb) ram
- 7b – 24gb ram without any background process
- 1.5b – 14gb ram

Trained llms

| NAME | SIZE |
|---------------------------------------|--------|
| • qwen2.5-coder:1.5b | 986 MB |
| • qwen-fuzzer-1.5-709-examples:latest | 3.1 GB |
| • qwen-fuzzer-1.5-172-examples:latest | 3.1 GB |

Conclusion

Key Achievements:

- Successfully demonstrated LLM-based fuzzing integration in CI/CD pipelines
- Achieved significant performance improvements over traditional methods
- Developed cost-effective solution using LoRA fine-tuning
- Created practical implementation within automotive constraints

Technical Contributions:

- Novel AI-assisted white-box fuzzing for automotive platforms
- Seamless CI/CD/CT pipeline integration with zero-downtime deployment
- Automated test artifact generation using fine-tuned LLMs
- Comprehensive evaluation framework with multiple metrics


Impact & Validation:

- Real-world validation in CARIAD's development environment
- Demonstrated cost reduction from thousands to under \$1 per campaign
- Enhanced security testing without impacting development velocity
- Established foundation for automotive-specific AI fuzzing frameworks

Future Research Directions:

- Integration with safety standards (ISO 26262, ISO 21434)
- Real-time constraint handling for automotive requirements
- Multi-ECU testing approaches for distributed architectures
- Standardization of safety-critical evaluation frameworks

„This thesis focuses on integrating AI and LLMs into CI/CD/CT pipelines to improve the security testing of automotive software.“



Vielen Dank
für Ihre Aufmerksamkeit!



QUESTIONS?